

Scheduling Many-Task Applications on Multi-Clouds and Hybrid Clouds

Shifat P. Mithila, Peter Franz, Gerald Baumgartner

Who am I?

- Undergraduate research assistant
- Upgraded the platform to multiple cloud sites
- Learned algorithms through working with the code and explanations

Status Quo Multi-Computer Approaches

Approach	Performance	Networking	Resources	Parallelism
Supercomputers	Predictable	Low Latency + Robust	Homogeneous	Fine Grained Parallelism
Desktop Grids	Unpredictable	High Latency + Unreliable	Heterogeneous	No parallelism; Independent Tasks
Cloud Virtual Resources	Unpredictable	Moderate Latency + Relialiable	Heterogeneous	Coarse Parallelism; Map Reduce

Virtual Machine Challenges

- Performance information is unpredictable before resources are provisioned
- Latency to neighbors varies as host pcs change
- Heterogeneous resources
- Performance variability - competition for host os resources

Supercomputer

- Forehand performance knowledge
- Robust and low latency networking
- Expensive
- Fine grained parallelism

Large-scale desktop-based master-worker grids

- Heterogenous and variable performance of nodes
- Unreliable and high latency networking
- Donated CPU cycles
- Little to no parallelism

Cloud Computing

- Virtual Machines
- Heterogenous and variable performance of nodes
- Robust but moderate latency networking
- Pay per use
- Fine grained parallelism (?)



Why you should be interested?

- Cheaper
- Higher availability
- Provider agnostic; allows price shopping
- Utilize your existing hardware
- Scale resources depending on task

A task based decentralized-vector-scheduling
platform, minimizing propagation delay with
latency based clustering.

Task Based

- Tasks break large problems into smaller steps
- Some tasks require other tasks to be completed to start
- Tasks are python script to run with passed parameters

$$\begin{aligned}
 R_{h_1 h_2}^{p_1 p_2} = & v_{vvoo}^{p_1 p_2}_{h_1 h_2} - t_{vo}^{p_1}_{h_7} * (v_{ovoo}^{h_7 p_2}_{h_1 h_2} - 0.5 * t_{vo}^{p_2}_{h_8} * \\
 & (v_{oooo}^{h_7 h_8}_{h_1 h_2} + (v_{ooov}^{h_7 h_8}_{h_1 p_3} - 0.5 * v_{oovv}^{h_7 h_8}_{p_3 p_4} * t_{vo}^{p_4}_{h_1}) * t_{vo}^{p_3}_{h_2} + \\
 & 0.5 * t_{vvoo}^{p_5 p_6}_{h_1 h_2} * v_{oovv}^{h_7 h_8}_{p_5 p_6}) + t_{vo}^{p_3}_{h_2} * (v_{ovov}^{h_7 p_2}_{h_1 p_3} - 0.5 * \\
 & t_{vo}^{p_4}_{h_1} * v_{ovvv}^{h_7 p_2}_{p_3 p_4}) - (f_{ov}^{h_7}_{p_3} - t_{vo}^{p_4}_{h_4} * v_{oovv}^{h_4 h_7}_{p_3 p_4}) * t_{vvoo}^{p_2 p_3}_{h_1 h_2} - \\
 & t_{vvoo}^{p_2 p_7}_{h_2 h_4} * (v_{ooov}^{h_4 h_7}_{h_1 p_7} + t_{vo}^{p_3}_{h_1} * v_{oovv}^{h_4 h_7}_{p_3 p_7}) + 0.5 * t_{vvoo}^{p_3 p_4}_{h_1 h_2} * \\
 & v_{ovvv}^{h_7 p_2}_{p_3 p_4} + t_{vo}^{p_3}_{h_2} * (v_{vvov}^{p_1 p_2}_{h_1 p_3} - 0.5 * t_{vo}^{p_4}_{h_1} * v_{vvvv}^{p_1 p_2}_{p_3 p_4}) + \\
 & (f_{oo}^{h_6}_{h_1} + (f_{ov}^{h_6}_{p_6} + t_{vo}^{p_4}_{h_4} * v_{oovv}^{h_4 h_6}_{p_4 p_6}) * t_{vo}^{p_6}_{h_1} - t_{vo}^{p_4}_{h_4} * \\
 & v_{ooov}^{h_4 h_6}_{h_1 p_4} - 0.5 * t_{vvoo}^{p_4 p_5}_{h_1 h_5} * v_{oovv}^{h_5 h_6}_{p_4 p_5}) * t_{vvoo}^{p_1 p_2}_{h_2 h_6} + (f_{vv}^{p_2}_{p_3} - \\
 & t_{vo}^{p_4}_{h_4} * v_{ovvv}^{h_4 p_2}_{p_3 p_4} - 0.5 * t_{vvoo}^{p_2 p_4}_{h_4 h_5} * v_{oovv}^{h_4 h_5}_{p_3 p_4}) * t_{vvoo}^{p_1 p_3}_{h_1 h_2} + \\
 & 0.5 * t_{vvoo}^{p_1 p_2}_{h_6 h_8} * (v_{oooo}^{h_6 h_8}_{h_1 h_2} + t_{vo}^{p_6}_{h_2} * (v_{ooov}^{h_6 h_8}_{h_1 p_6} + 0.5 * \\
 & t_{vo}^{p_4}_{h_1} * v_{oovv}^{h_6 h_8}_{p_4 p_6}) - 0.5 * t_{vvoo}^{p_3 p_4}_{h_1 h_2} * v_{oovv}^{h_6 h_8}_{p_3 p_4}) + t_{vvoo}^{p_1 p_3}_{h_2 h_3} * \\
 & (v_{ovov}^{h_3 p_2}_{h_1 p_3} - t_{vo}^{p_5}_{h_1} * v_{ovvv}^{h_3 p_2}_{p_3 p_5} - 0.5 * t_{vvoo}^{p_2 p_5}_{h_1 h_5} * v_{oovv}^{h_3 h_5}_{p_3 p_5}) + \\
 & 0.5 * t_{vvoo}^{p_3 p_4}_{h_1 h_2} * v_{vvvv}^{p_1 p_2}_{p_3 p_4}
 \end{aligned}$$

Spin-orbital coupled-cluster singles-and-double doubles
equation

Task Passing


- All nodes must be utilized
- Nodes must share
- Must be a fast decision
- Balancing act between multiple considerations

Decentralized


- Avoiding a centralized bottleneck
- Short-running parallel-tasks on network connected heterogeneous resources
 - Controller is remote from nodes in cloud environment
 - Short running means less time for decisions to come from the controller
 - Controller needs to manage large amounts performance information
- Need task passing decision to be made at node level


Vector-Scheduling

- Lightweight decentralized algorithm
- Nodes connected in overlay graph
- Nodes send performance information to neighbors
- Measurements are normalized
- Using an experiment defined flow vector tasks sent in desired direction to neighbors




Vector Scheduling worked example





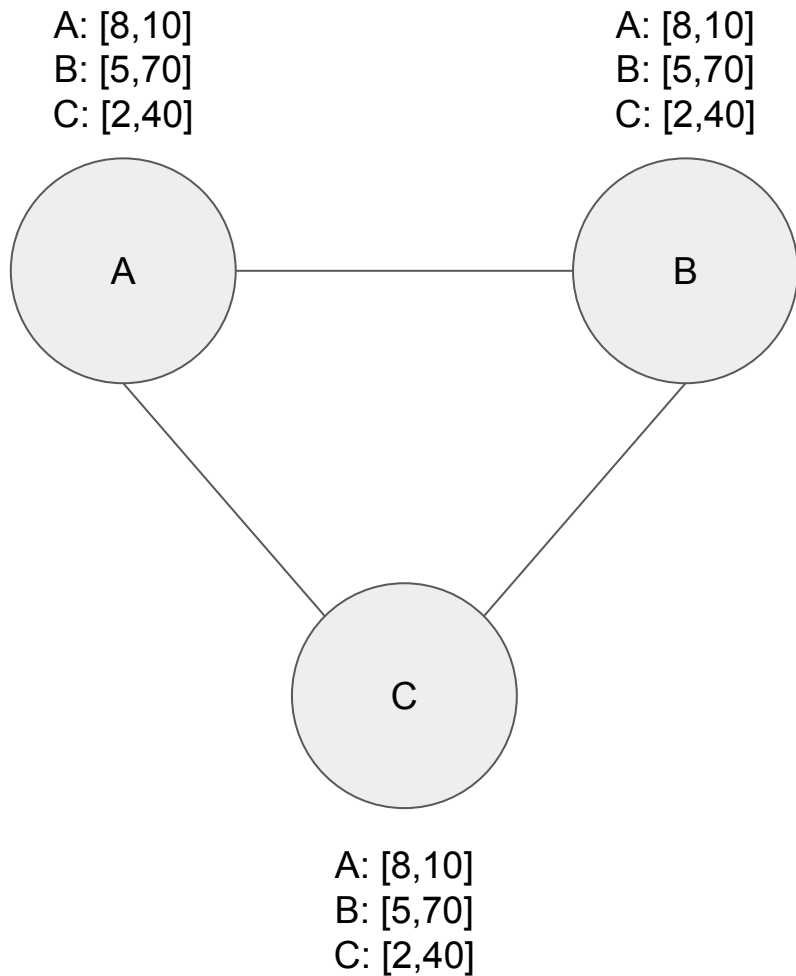
Nodes probe a characteristic

```
Ql = self.queueLen()  
Perf = self.runBenchmark()
```



Nodes provide information to neighbors


For each neighbor:
shipQLToNeighbor(neighbor)
shipPerfToNeighbor(neighbor)



Nodes keep track of the components in
a vector for each neighbor
[QueueLength, Performance]

```
For each neighbor:  
  neighborVector =  
  [  
    reciveQL(neighbor),  
    receivePerf(neighbor),  
  ]
```

A: [1,-1]
B: [0,1]
C: [-1,0]




A: [1,-1]
B: [0,1]
C: [-1,0]

A: [1,-1]
B: [0,1]
C: [-1,0]

Nodes normalize vectors

Scoring with a flow-vector



Flow Vector:
[-1,-0.25]

Node	Queue Length	Performance	Score
A	1	-1	-0.75
B	0	1	-0.25
C	-1	0	1

$$A: (1 * -1) + (-1 * -0.25) = -0.75$$


$$B: (0 * -1) + (1 * -0.25) = -0.25$$

$$C: (-1 * -1) + (0 * -0.25) = 1$$

Thrashing

- Tasks passed back and forth with no completion
- Want passing, but not for minor gains
- “Statis vector” to minimize thrashing

Stasis-Vector



Flow Vector:
[-1,-0.25]

Stasis Vector:
[0,-0.5]

Node	Queue Length	Performance	Score
A	1	-1	-0.75 -0.625
B	0	1	-0.25
C	-1	0	1

$$A: ([1+0] * -1) + ([-1 + -0.5] * -0.25) = -0.625$$

Adding latency


- Vector scheduling is an extremely expandable
- Add Latency to scheduling
- Allows us to consider nodes relative locations

Network Connectivity

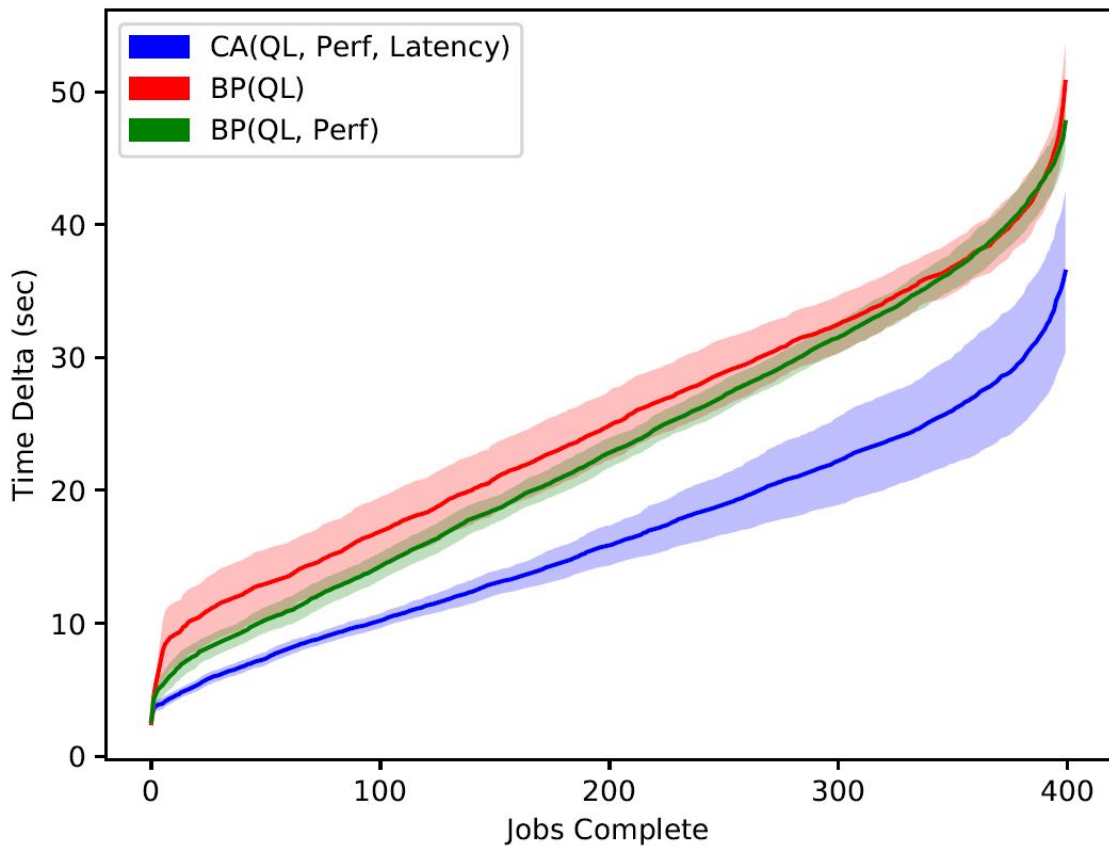
- Full connectivity between all nodes has poor performance
- High latency connections should be avoided
- More neighbors increases overhead of vector scheduling

Centralized Clustering Algorithm

- Measure latency to neighbors
 - Nodes form a full graph with closest neighbors, Clique
 - Keep some connections to distant groups
-
- Reduces number of neighbors to evaluate
 - Removes slower connections
 - Maintains connectivity




Latency + Clustering increases performance



Hybrid + Multi-Site Resource Proof of Concept

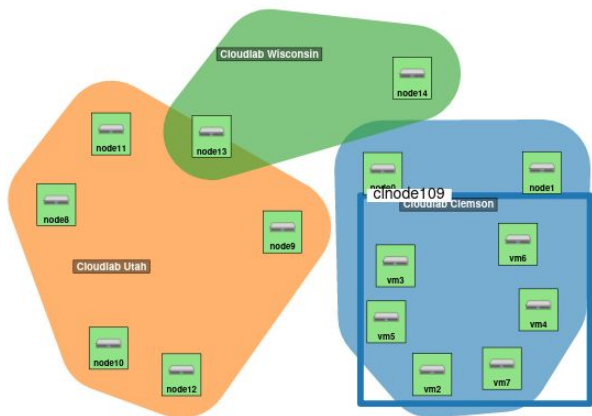
- 400 65 x 65 matrix multiplication tasks, with 50KB payload file included
- Demonstrates our capacity to stitch together multiple resource types and sites
- Limitation with handling LSU firewall, more development needed

experiment	resources	initial vector	final vector	time
Local	2 PCs	$[-1, -0.5, 0.5]$	$[-1, 0, -0.25]$	457.95
Hybrid	2 PCs, 2 VMS	$[-1, -0.5, 0.5]$	$[-1, 0, -0.25]$	280.47



Extending Research with Heterogeneity

- 400 65 x 65 matrix multiplication tasks, with 50KB payload file
- Heterogeneous performance
- Random graph construction, 20%



Site	Hardware Type	Workers	CPU	CPU Speed MHZ
Clemson	VM	4	2 cores per	-
Clemson	c4130	10	Intel E5-2680v3	2500
Utah	d6515	10	AMD EPYC 7452	2400
Utah	m510	12	Intel Xeon D-1548	2000
Wisconsin	c220g2	10	Haswell	2600

3D Vector Scheduling in Multi-Site Cloud

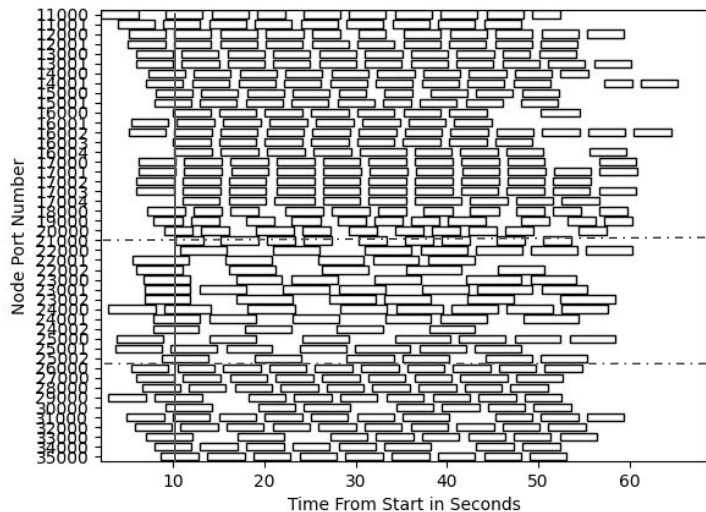
- Varying vector configurations
- Ignoring latency had worst performance
- Avoiding latency is good for the middle and end
- Intentionally pushing tasks at the start, leads to good task saturation

Table 1: Performance comparison for different [queue, performance, latency] scheduling vectors with vector swap after 30 seconds.

Experiment	initial vector	final vector	time
Ignore latency	$[-1, -0.5, 0]$	$[-1, 0.5, 0]$	78.9
Emphasize queue	$[-1, 0, 0]$	$[-0.5, -0.5, -0.25]$	72.9
Emphasize performance	$[-1, -0.5, 0.5]$	$[-1, 0, -0.25]$	67.6


Reading Tea Leaves

Use Latency + Push Away to Start



- Using latency to start has 93% Node utilization after 10 seconds
- All sites have work until the end

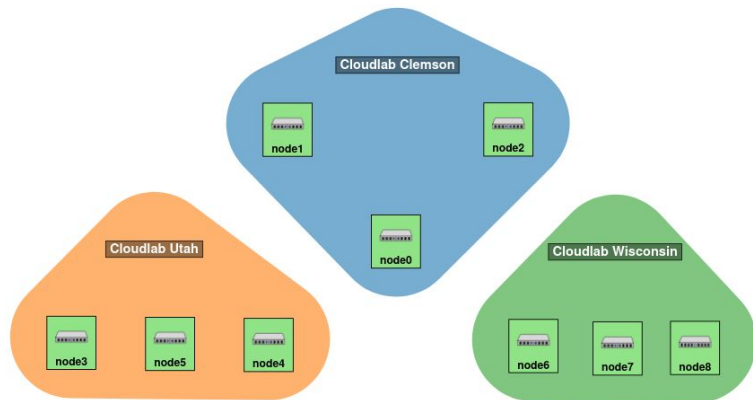
Ignore Latency



- Ignoring latency has 73% Node utilization after 10 seconds
- Some sites starve at the end

Extending Research with Heterogeneity

- 400 65 x 65 matrix multiplication tasks, with 50KB payload file
- Heterogeneous performance
- Random graph construction, 20%

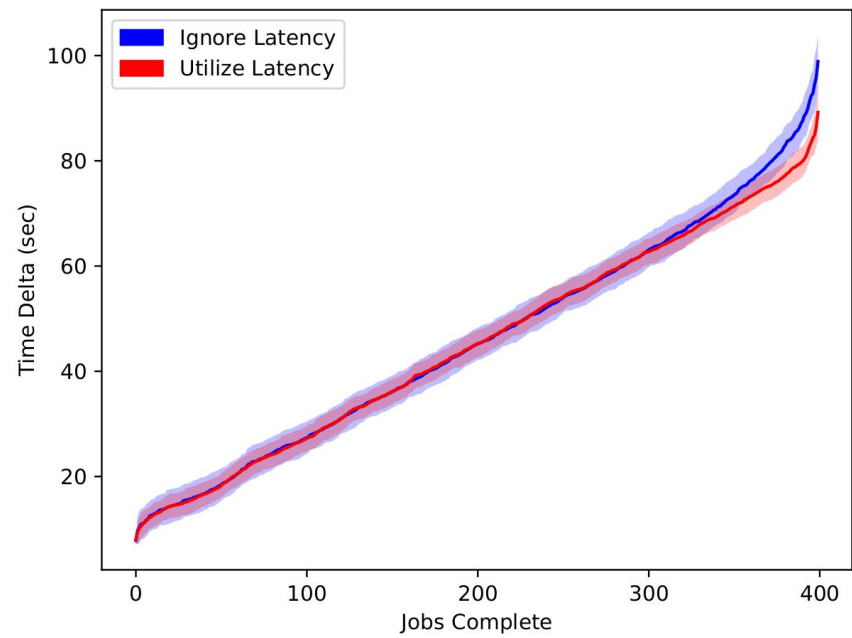


Site	Hardware Type	Workers	CPU	CPU Speed MHZ
Clemson	c8220	10	Intel E5-2660v2	2200
Utah	m510	16	Intel Xeon D-1548	2000
Wisconsin	c220g1	20	Haswell	2400

3D Vector Scheduling in Multi-Site Cloud


- Varying vector configurations
- Ignoring latency had worst performance
- Intentionally pushing tasks at the start, leads to good task saturation

Name	Initial vector	Second Vector	Average Time
Ignore Latency	-1,-0.5,0	-1,0.5,0	100.93
Minimize Latency	-1,-0.25,-0.25	-1,0,-0.25	91.55
Push Tasks	-1,-0.5,0.5	-1,0,-0.25	89.82




Reading Tea Leaves

Use Latency + Push Away to Start



- Nodes quickly utilized
- All sites have work until the end

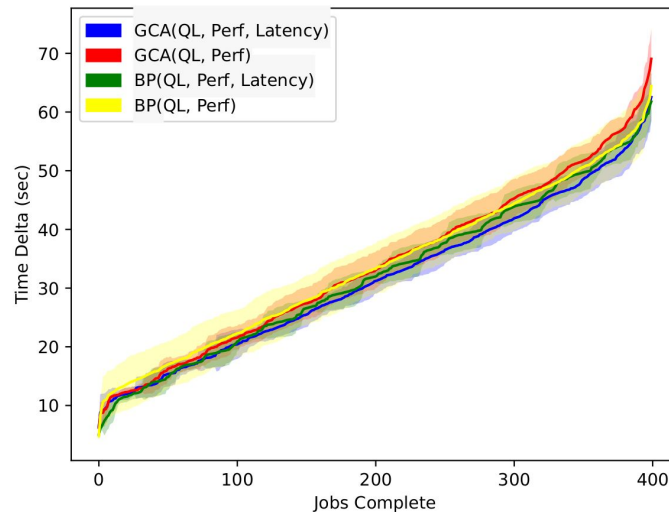
Ignore Latency



- Slow to utilize all nodes
- Some sites starve at the end

Clustering Algorithm Limitations

- 46 worker nodes, 400 65x65 Matrix Mult tasks
 - GCA = Graph Construction Algorithm, Cluster
 - BP = Brain Peterson Algorithm, 20% Connectivity
- Clustering algorithm on real world clusters have oversize cliques
- Limited trials



Further Work

- Improve adaptation of clustering algorithm parameters
- Trial automation
- General algorithm for changing flow vectors at runtime
- Probe hardware and system information to include as a vector component
- Task dependencies
- Tasks ran on collaboration of multiple nodes