

Lecture 21: Domain-Specific Programming Systems

Parallel Computer Architecture and Programming

CMU 15-418/15-618, Spring 2023

What we have learnt

- Design computer systems that can **scale**
 - Running faster given more resources
- Design computer systems that are **efficient**
 - Running faster under resource constraints
- Techniques discussed
 - Exploiting **parallelism** in applications
 - Exploiting **locality** in applications
 - Leveraging hardware **specialization**

Various programming models to abstract hardware

Machines with very different performance characteristics

- CPUs, GPUs, TPUs, systolic arrays

Different technologies and performance characteristics within the same machine at different scales

- **Within a core**: SIMD, multi-threading: fine grained sync and comm
 - Abstractions: SPMD programming (ISPC, CUDA, OpenCL)
- **Across cores**: coherent shared memory via fast on-chip network
 - Abstractions: OpenMP pragma, Cilk
- **Across racks**: distributed memory, multi-stage network
 - Abstractions: message passing (MPI, Go, Spark, Legion, Charm++)

Various programming models to abstract hardware

Machines with very different performance characteristics

- CPUs, GPUs, TPUs, systolic arrays

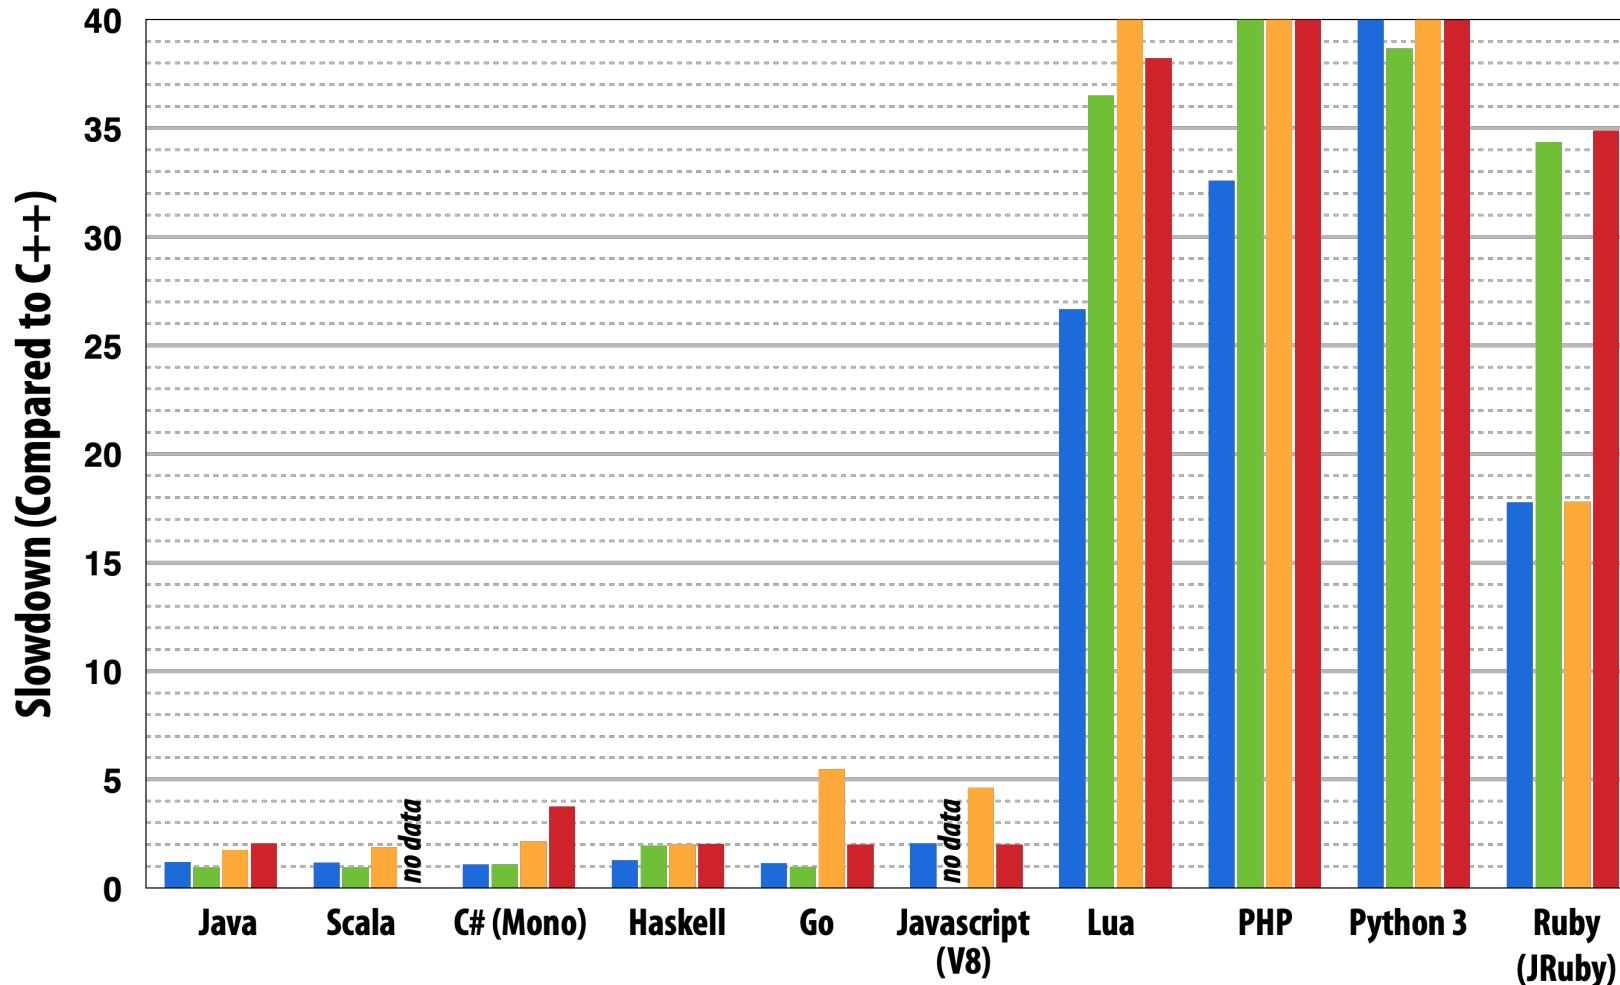
Different technologies and performance characteristics within the same machine at different scales

To be efficient, software must be optimized for HW characteristics

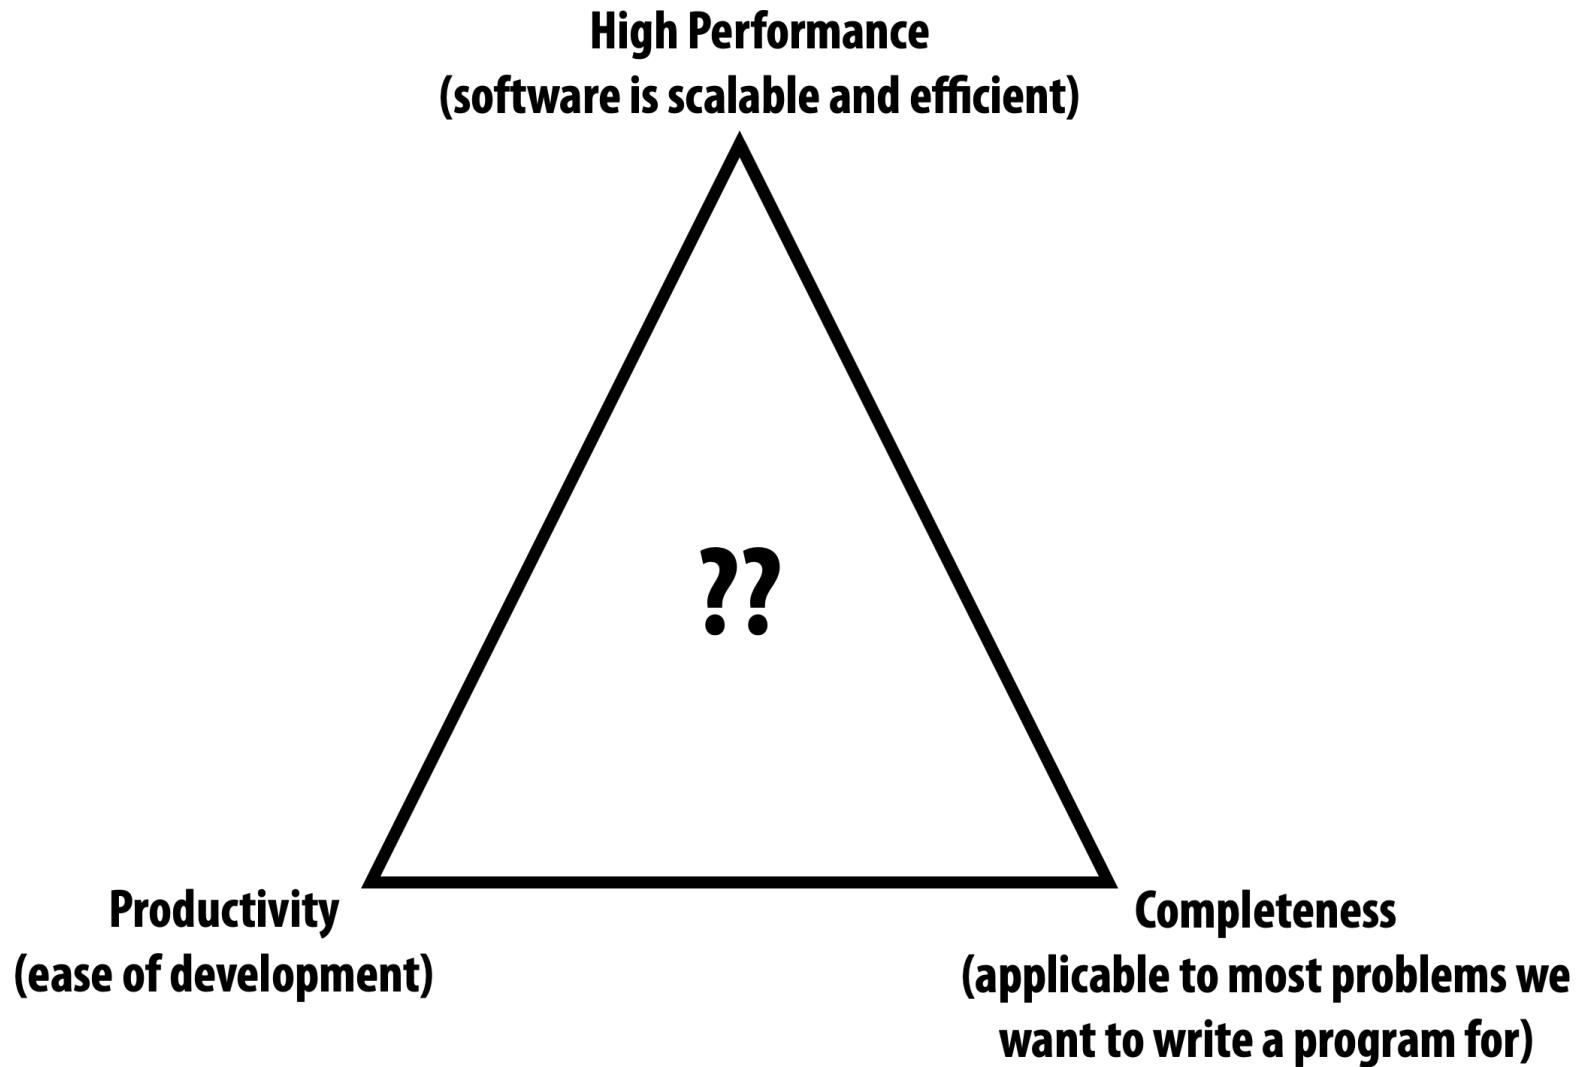
- Difficult even in the case of one level of one machine
- **Combinatorial complexity of optimizations** when considering a complex machine, or different machines
- Loss of **software portability**

Most software systems use hardware inefficiently

Compared against GCC -o3 (no manual vector optimizations)

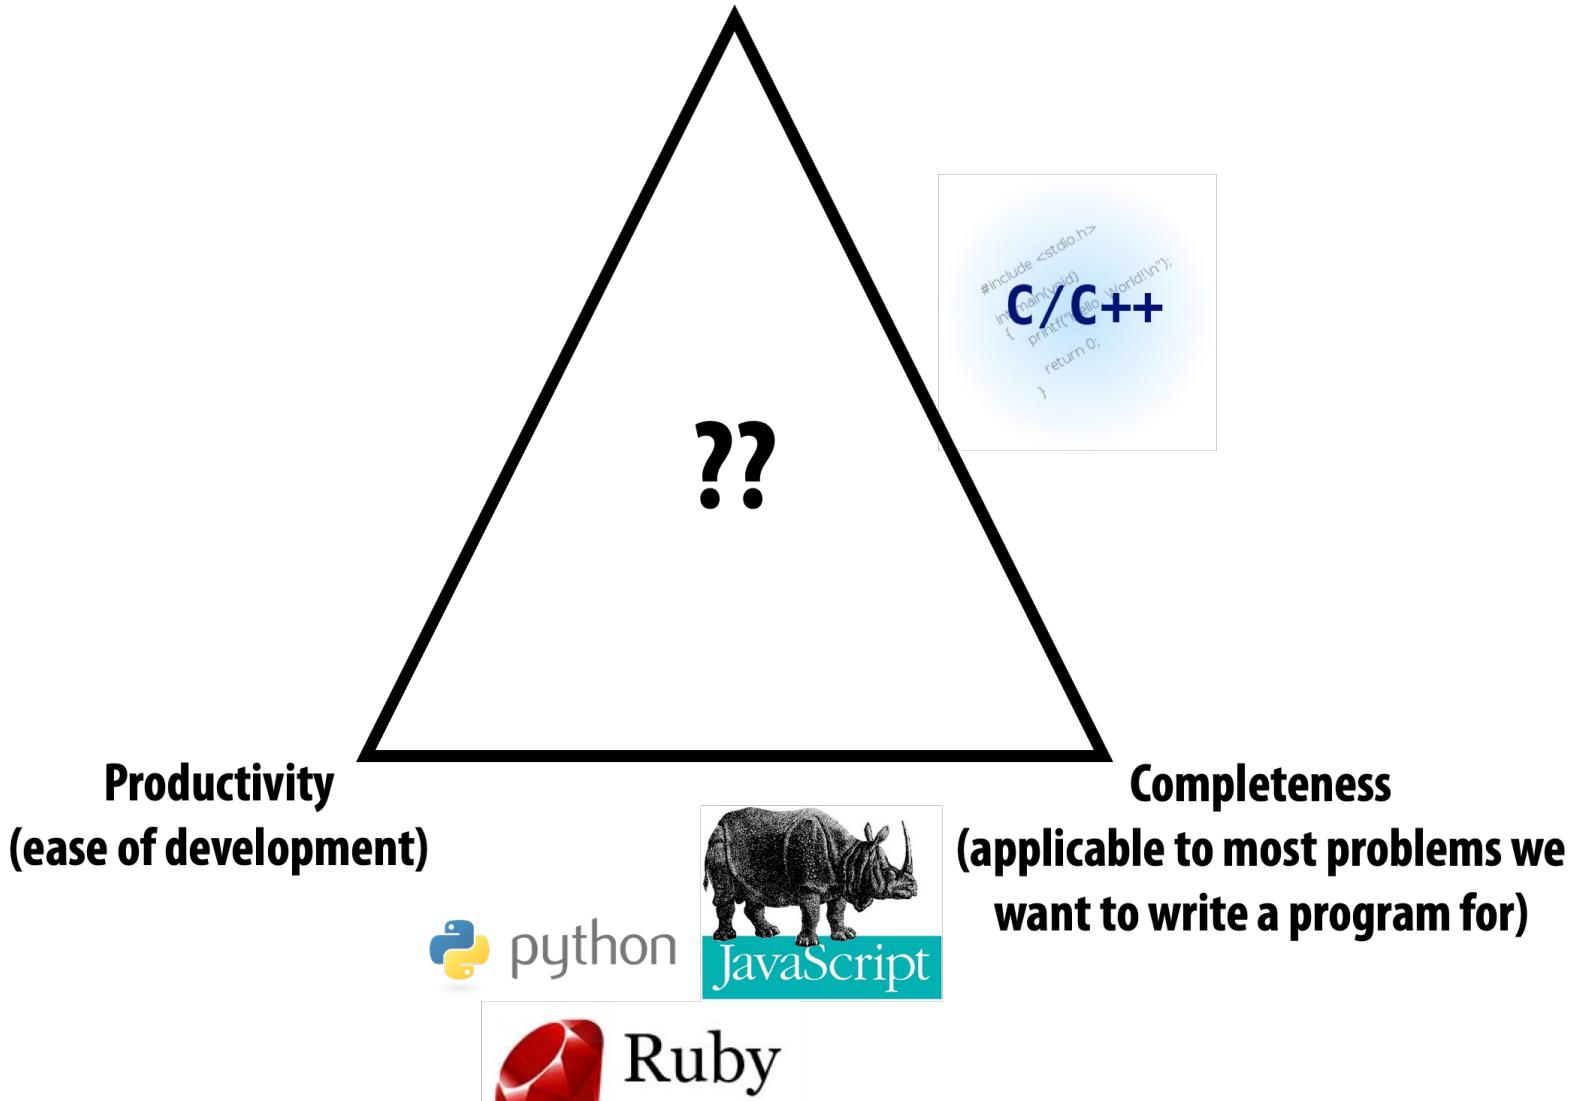


The Magical Ideal Parallel Programming System

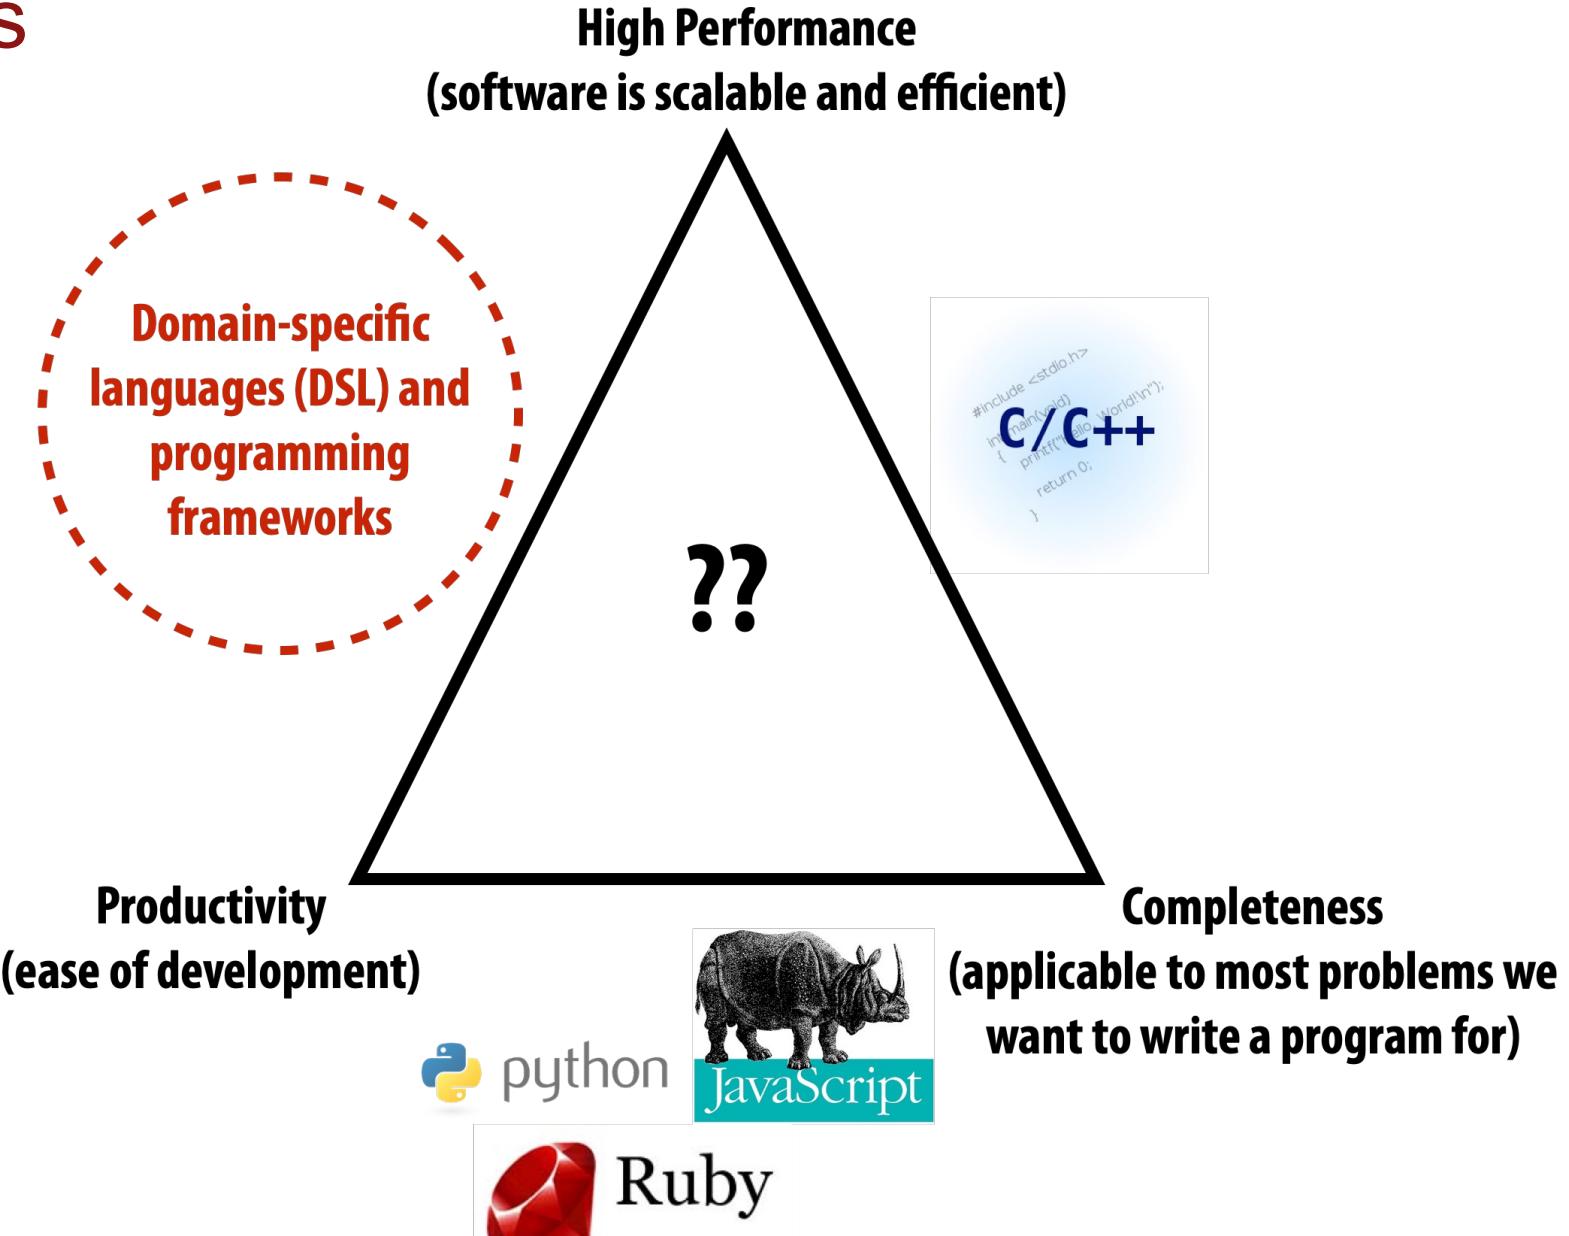


Widely Used Programming Languages

High Performance
(software is scalable and efficient)



Growing Interest in Domain-Specific Programming Systems



Domain-Specific Programming Systems

Key idea: raise level of abstraction for expressing programs

Introduce high-level programming primitives specific to an application domain

- **Productive**: intuitive to use, portable across machines, primitives correspond to behaviors frequently used to solve problems in targeted domain
- **Performant**: system uses domain knowledge to provide efficient, optimized implementation(s)
 - Given a machine: system knows what algorithms to use, parallelization strategies to employ for this domain

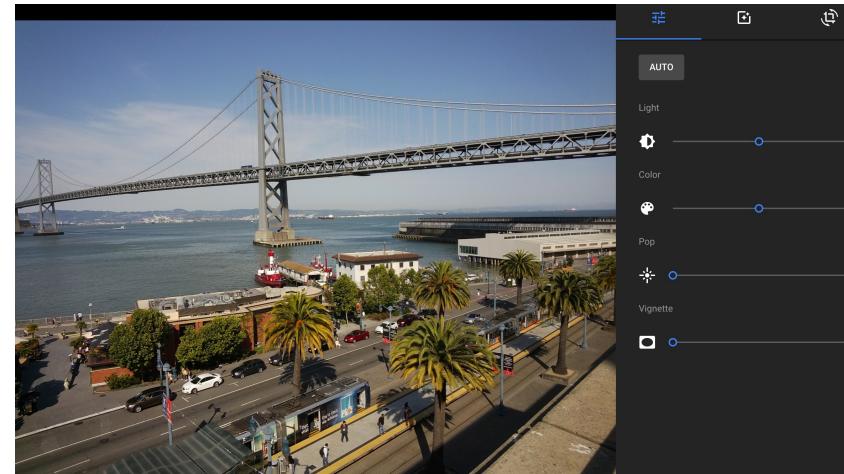
Cost: loss of generality/completeness

Two Domain-Specific Programming Systems

1. Halide: for image processing
2. TVM: for deep learning

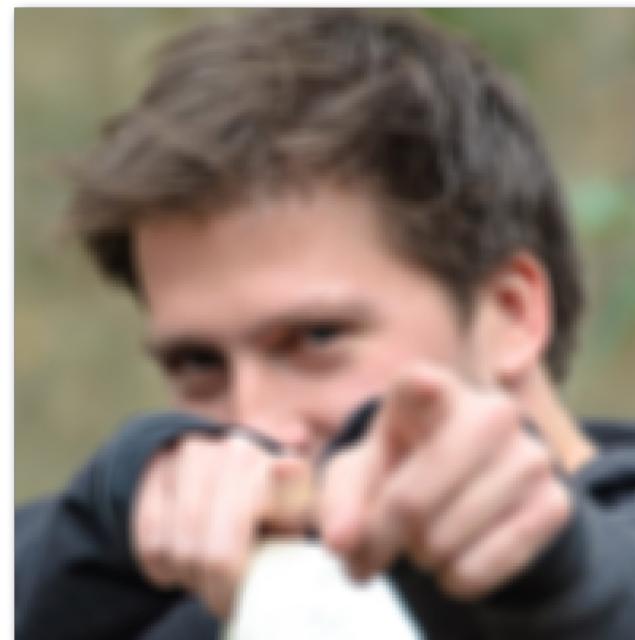
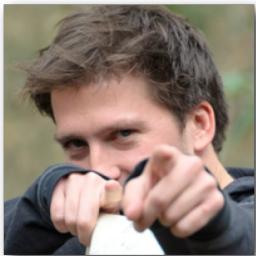
Halide: a Domain-Specific Language for Image Processing

- Used to implemented Android HDR+ app
- Halide code used to process all images uploaded to Google Photos



A Quick Tutorial on High-Performance Image Processing

Image Blur



(Zoom view)

3x3 Image Blur (a convolution with predefined weights)

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/9, 1.0/9, 1.0/9,
                   1.0/9, 1.0/9, 1.0/9,
                   1.0/9, 1.0/9, 1.0/9};

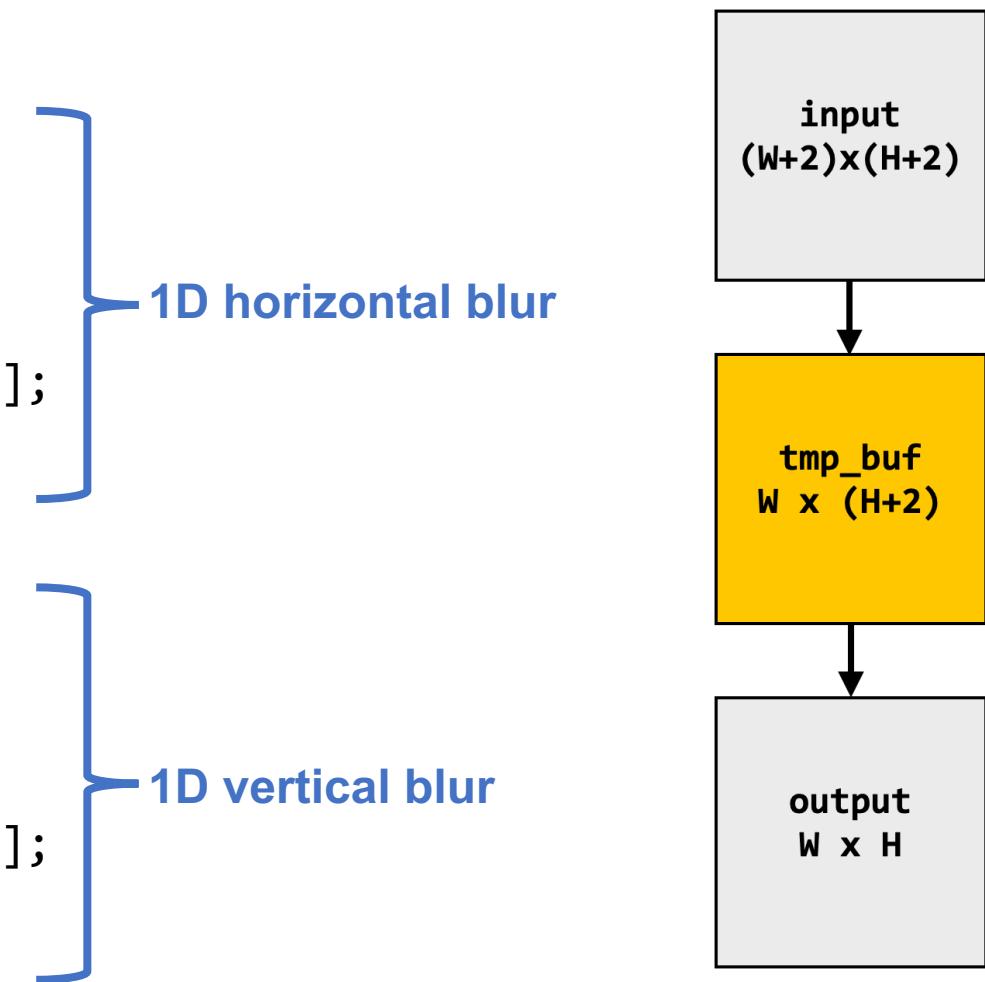
for (int j=0; j<HEIGHT; j++) {
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int jj=0; jj<3; jj++)
            for (int ii=0; ii<3; ii++)
                tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];
        output[j*WIDTH + i] = tmp;
    }
}
```

Total work per image: $9 * \text{WIDTH} * \text{HEIGHT}$
For NxN filter: $N * N * \text{WIDTH} * \text{HEIGHT}$

Two-Pass 3x3 Blur

```
float input[(WIDTH+2) * (HEIGHT+2)];  
float tmp_buf[WIDTH * (HEIGHT+2)];  
float output[WIDTH * HEIGHT];  
  
float weights[] = {1.0/3, 1.0/3, 1.0/3};  
for (int j=0; j<(HEIGHT+2); j++)  
    for (int i=0; i<WIDTH; i++) {  
        float tmp = 0.f;  
        for (int ii=0; ii<3; ii++)  
            tmp += input[j*(WIDTH+2) + i+ii] * weights[ii];  
        tmp_buf[j*WIDTH + i] = tmp;  
    }  
  
for (int j=0; j<HEIGHT; j++)  
    for (int i=0; i<WIDTH; i++) {  
        float tmp = 0.f;  
        for (int jj=0; jj<3; jj++)  
            tmp += tmp_buf[(j+jj)*WIDTH + i] * weights[jj];  
        output[j*WIDTH + i] = tmp;  
    }
```

Total work per image: $6 * \text{WIDTH} * \text{HEIGHT}$
For NxN filter: $2 * N * \text{WIDTH} * \text{HEIGHT}$
Extra memory: $\text{WEIGHT} * \text{HEIGHT}$
3x lower arithmetic intensity



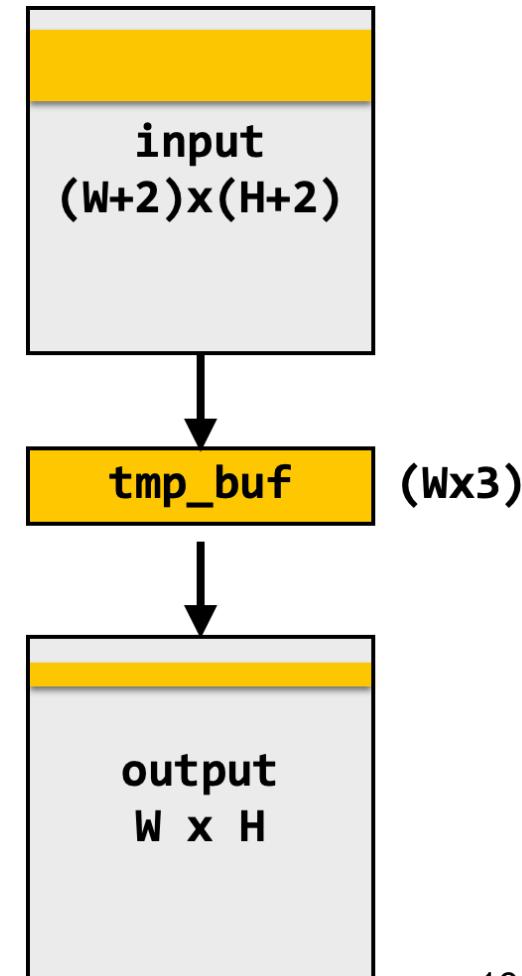
Two-Pass 3x3 Blur (Chunked)

```
float input[(WIDTH+2) * (HEIGHT+2)];  
float tmp_buf[WIDTH * 3];  
float output[WIDTH * HEIGHT];  
float weights[] = {1.0/3, 1.0/3, 1.0/3};  
  
for (int j=0; j<HEIGHT; j++) {  
    for (int j2=0; j2<3; j2++)  
        for (int i=0; i<WIDTH; i++) {  
            float tmp = 0.f;  
            for (int ii=0; ii<3; ii++)  
                tmp += input[(j+j2)*(WIDTH+2) + i+ii] * weights[ii];  
            tmp_buf[j2*WIDTH + i] = tmp; ————— Produce 3 rows of tmp_buf  
  
    for (int i=0; i<WIDTH; i++) {  
        float tmp = 0.f;  
        for (int jj=0; jj<3; jj++)  
            tmp += tmp_buf[jj*WIDTH + i] * weights[jj];  
        output[j*WIDTH + i] = tmp;  
    }  
}
```

Total work per row of output:

- step 1: $3 \times 3 \times \text{WIDTH}$ work
- step 2: $3 \times \text{WIDTH}$ work

Total work per image = $12 \times \text{WIDTH} \times \text{HEIGHT}$
Loads from tmp_buffer are cached



Combine them together
to get one row of output

Conflicting goals (once again...)

- Want to be computationally efficient (perform fewer operations)
- Want to take advantage of locality when possible
 - Otherwise computationally efficient code will be bandwidth bound
- Want to execute in parallel (multi-core, SIMD within core)

Optimized C++ code: 3x3 image blur

```
void fast_blur(const Image &in, Image &blurred) {
    _m128i one_third = _mm_set1_epi16(21846);
#pragma omp parallel for
for (int yTile = 0; yTile < in.height(); yTile += 32) {
    _m128i a, b, c, sum, avg;
    _m128i tmp[(256/8)*(32+2)];
    for (int xTile = 0; xTile < in.width(); xTile += 256) {
        _m128i *tmpPtr = tmp;
        for (int y = -1; y < 32+1; y++) {
            const uint16_t *inPtr = &(in(xTile, yTile+y));
            for (int x = 0; x < 256; x += 8) {
                a = _mm_loadu_si128((__m128i*) (inPtr-1));
                b = _mm_loadu_si128((__m128i*) (inPtr+1));
                c = _mm_load_si128((__m128i*) (inPtr));
                sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                avg = _mm_mulhi_epi16(sum, one_third);
                _mm_store_si128(tmpPtr++, avg);
                inPtr += 8;
            }
            tmpPtr = tmp;
            for (int y = 0; y < 32; y++) {
                _m128i *outPtr = ((__m128i *) )(&(blurred(xTile, yTile+y)));
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_load_si128(tmpPtr+(2*256)/8);
                    b = _mm_load_si128(tmpPtr+256/8);
                    c = _mm_load_si128(tmpPtr++);
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulhi_epi16(sum, one_third);
                    _mm_store_si128(outPtr++, avg);
                }
            }
        }
    }
}
}}}}
```

- + 10x faster than the original code
- Specific to SSE (not AVX2), CPU-code only
- Lacks readability, portability, modularity

Halide: Decouple Algorithm from Schedule

- Algorithm: *what* to do
- Schedule: *how* to do

```
Func halide_blur(Func in) {
    Func tmp, blurred;
    Var x, y, xi, yi;

    // The algorithm
    tmp(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;
    blurred(x, y) = (tmp(x, y-1) + tmp(x, y) + tmp(x, y+1))/3;

    // The schedule
    blurred.tile(x, y, xi, yi, 256, 32)
        .vectorize(xi, 8).parallel(y);
    tmp.chunk(x).vectorize(x, 8);

    return blurred;
}
```

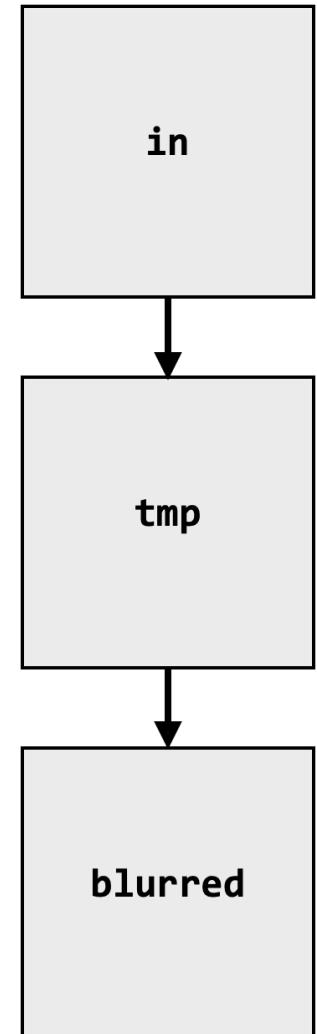
Why Decoupling Algorithm from Schedule?

- Algorithm: *what* to do
 - Schedule: *how* to do
-
- Easy for programmers to build pipelines
 - Easy for programmers to specify & explore optimizations
 - Easy for compilers to generate fast code

Algorithm: Pure Functional

- Declarative specification
- Pipeline stages are pure functions from coordinates to values
- No explicit bounds
- No loops or traversal orders
- Only feed forward pipelines

```
// The algorithm
tmp(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;
blurred(x, y) = (tmp(x, y-1) + tmp(x, y) + tmp(x, y+1))/3;
```

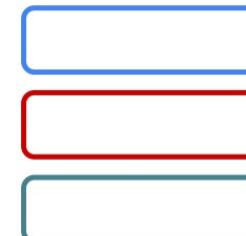


Think of a Halide algorithm as a pipeline

Schedule: describe how to execute a pipeline

- Defines intra-stage order and inter-stage interleaving
- For each stage:
 - 1) In which order should we compute its values?
 - 2) How to map onto parallel execution resources like SIMD units and GPU blocks?

```
// The schedule
blurred.tile(x, y, xi, yi, 256, 32)
    .vectorize(xi, 8).parallel(y);
tmp.chunk(x).vectorize(x, 8);
```

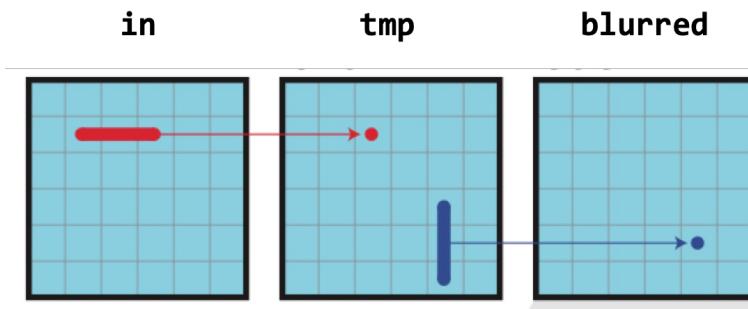


traversal order

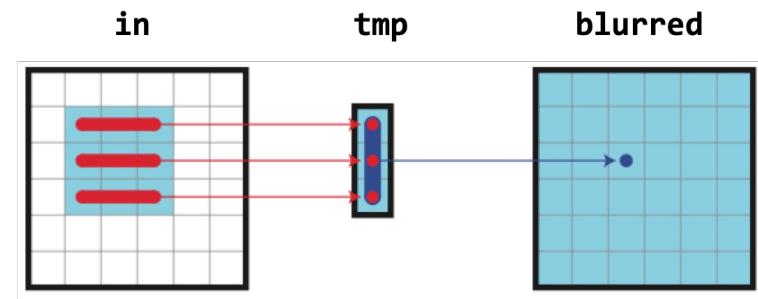
parallel execution

producer consumer relation

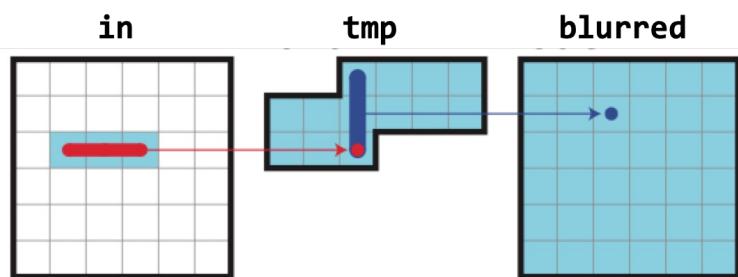
Producer/Consumer Scheduling Primitives



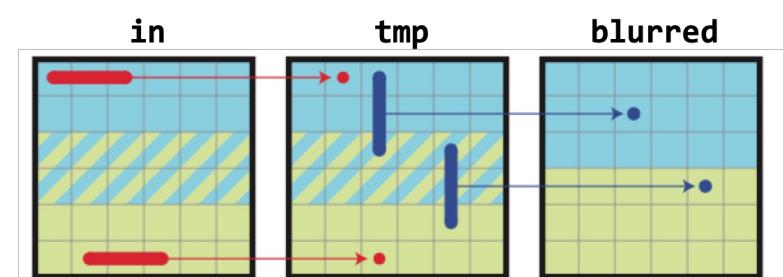
breadth first: each function is entirely evaluated before the next one.
"Root"



total fusion: values are computed on the fly each time that they are needed.
"Inline"



sliding window: values are computed when needed then stored until not useful anymore.
"Sliding Window"



tiles: overlapping regions are processed in parallel, functions are evaluated one after another.
"Chunked"

Producer/Consumer Scheduling Primitives

```
// Halide program definition
Func halide_blur(Func in) {

    Func blurx, out;
    Var x, y, xi, yi

    // the "algorithm description" (what to do)
    blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
    out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;

    // "the schedule" (how to do it)
    blurx.compute_at(ROOT);
    return out;
}
```

"Root":
compute all points of the producer,
then run consumer (minimal locality)

```
void halide_blur(uint8_t* in, uint8_t* out) {
    uint8_t blurx[WIDTH * HEIGHT];

    for (int y=0; y<HEIGHT; y++) {
        for (int x=0; y<WIDTH; x++) {
            blurx[] = ...

    for (int y=0; y<HEIGHT; y++) {
        for (int x=0; y<WIDTH; x++) {
            out[] = ...
    }
}
```

```
// Halide program definition
Func halide_blur(Func in) {

    Func blurx, out;
    Var x, y, xi, yi

    // the "algorithm description" (what to do)
    blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
    out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;

    // "the schedule" (how to do it)
    blurx.inline();
    return out;
}
```

"Inline":
revalue producer at every use site
in consumer (maximal locality)

```
void halide_blur(uint8_t* in, uint8_t* out) {
    for (int y=0; y<HEIGHT; y++) {
        for (int x=0; y<WIDTH; x++) {
            out[] = (((in[(y-1)*WIDTH+x-1] +
                in[(y-1)*WIDTH+x] +
                in[(y-1)*WIDTH+x+1]) / 3) +
                ((in[y*WIDTH+x-1] +
                in[y*WIDTH+x] +
                in[y*WIDTH+x+1]) / 3) +
                ((in[(y+1)*WIDTH+x-1] +
                in[(y+1)*WIDTH+x] +
                in[(y+1)*WIDTH+x+1]) / 3));
    }
}
```

Schedule: describe how to execute a pipeline

```
// Halide program definition
Func halide_blur(Func in) {

    Func blurx, out;
    Var x, y, xi, yi

    // the "algorithm description" (what to do)
    blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
    out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;

    // "the schedule" (how to do it)
    out.tile(x, y, xi, yi, 256, 32).vectorize(xi,8).parallel(y);
    blurx.chunk(x).vectorize(x, 8);
    return out;
}

void halide_blur(uint8_t* in, uint8_t* out) {
    #pragma omp parallel for
    for (int y=0; y<HEIGHT; y+=32) {          // tile loop
        for (int x=0; y<WIDTH; x+=256) {        // tile loop

            // buffer
            uint8_t* blurx[34 * 256];

            // produce intermediate buffer
            for (int yi=0; yi<34; yi++) {
                // SIMD vectorize this loop (not shown)
                for (int xi=0; xi<256; xi++) {
                    blurx[yi*256+xi] =
                        (in[(y+yi-1)*WIDTH+x+xi-1] +
                         in[(y+yi-1)*WIDTH+x+xi] +
                         in[(y+yi-1)*WIDTH+x+xi+1]) / 3.0;
                }
            }

            // consumer intermediate buffer
            for (int yi=0; yi<32; yi++) {
                // SIMD vectorize this loop (not shown)
                for (int xi=0; xi<256; xi++) {
                    out[(y+yi)*256+(x+xi)] =
                        (blurx[yi*256+xi] +
                         blurx[(yi+1)*256+xi] +
                         blurx[(yi+2)*256+xi]) / 3.0;
                }
            }
        } // loop over tiles
    } // loop over tiles
}
```

Given a schedule, Halide carries out mechanical process of implementing the specified schedule

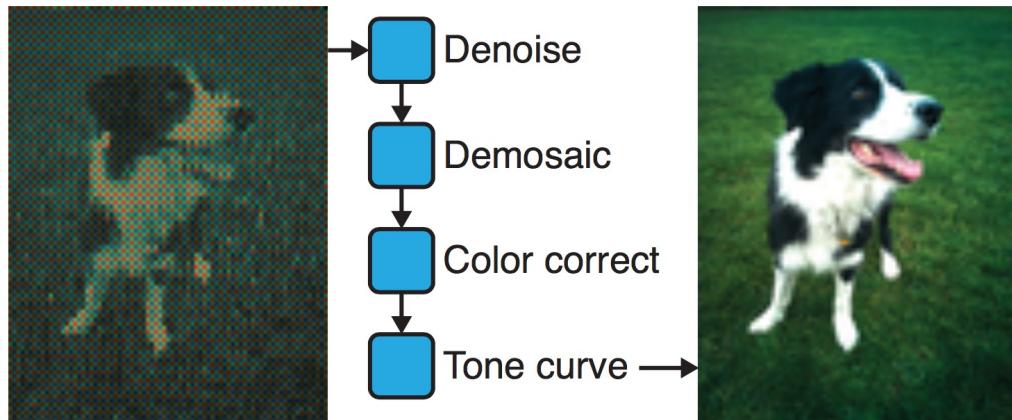
Halide: two domain-specific co-languages

- **Functional language** for describing image processing operations
- **Domain-specific language** for describing schedules
- **Design principle:** separate “algorithm specification” from its schedule
 - Programmer’s responsibility: provide a high-performance schedule
 - Compiler’s responsibility: carry out mechanical process of generating threads, SIMD instructions, managing buffers, etc.
 - Result: enable programmer to rapidly explore space of schedules
 - (e.g., “tile these loops”, “vectorize this loop”, “parallelize this loop across cores”)
- **Domain scope:**
 - All computation on regular N-D coordinate spaces
 - Only feed-forward pipelines
 - All dependencies inferable by compiler

Example Halide Results

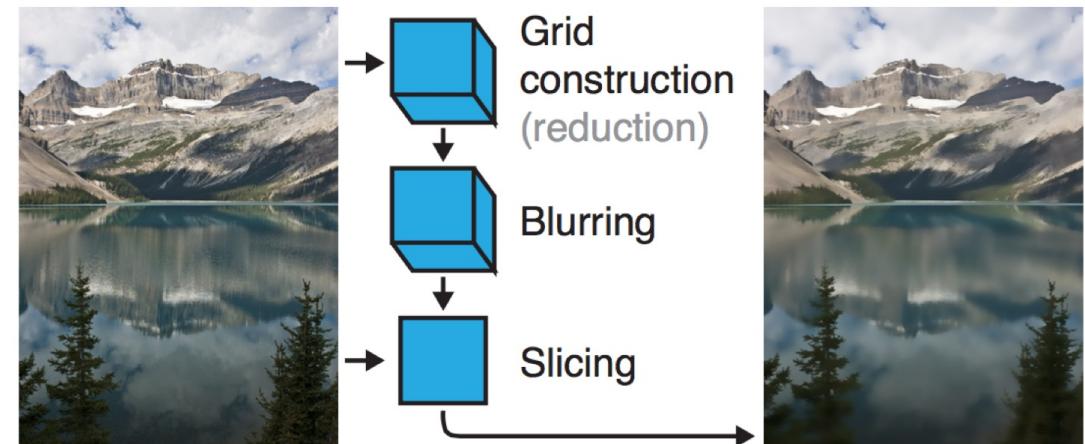
Camera RAW processing pipeline
(Convert RAW sensor data to RGB)

- Original: **463** lines of hand-tuned ARM NEON assembly
- Halide: **2.75x** less code, **5%** faster



Bilateral filter

- Original 122 lines of C++
- Halide: 34 lines algorithm + 6 lines schedule
 - CPU implementation: 5.9x faster
 - GPU implementation: 2x faster than hand-written CUDA



Recap: Halide is a DSL

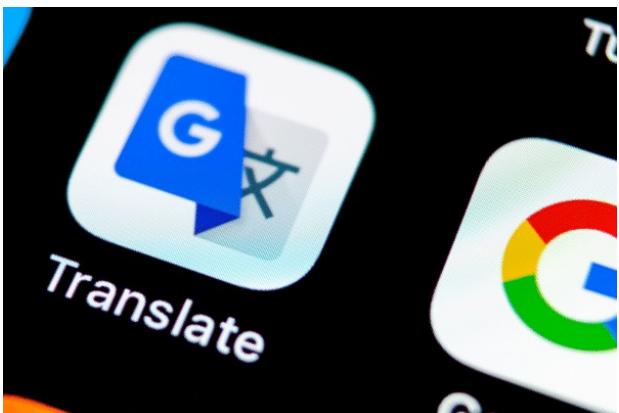
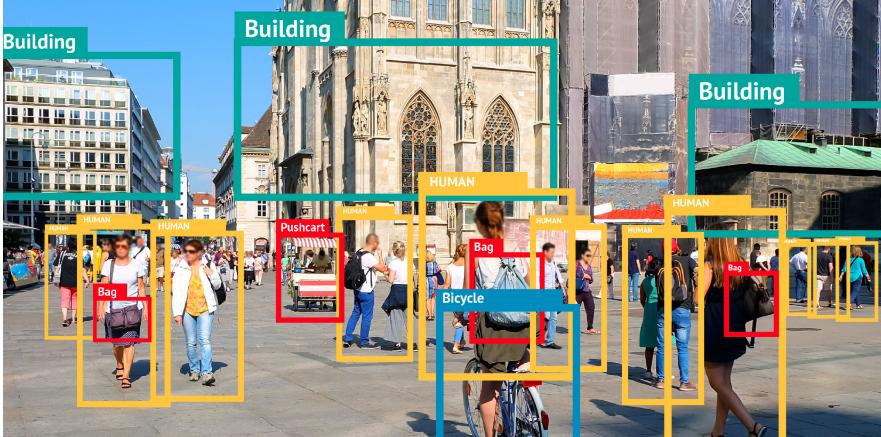
For helping developers optimize image processing code more rapidly

- Halide doesn't decide how to optimize a program for a novice programmer
- Halide provides primitives for a programmer to rapidly express what optimizations the system should apply
- **Halide carries out the nitty-gritty of mapping that strategy to a machine**

Two Domain-Specific Programming Systems

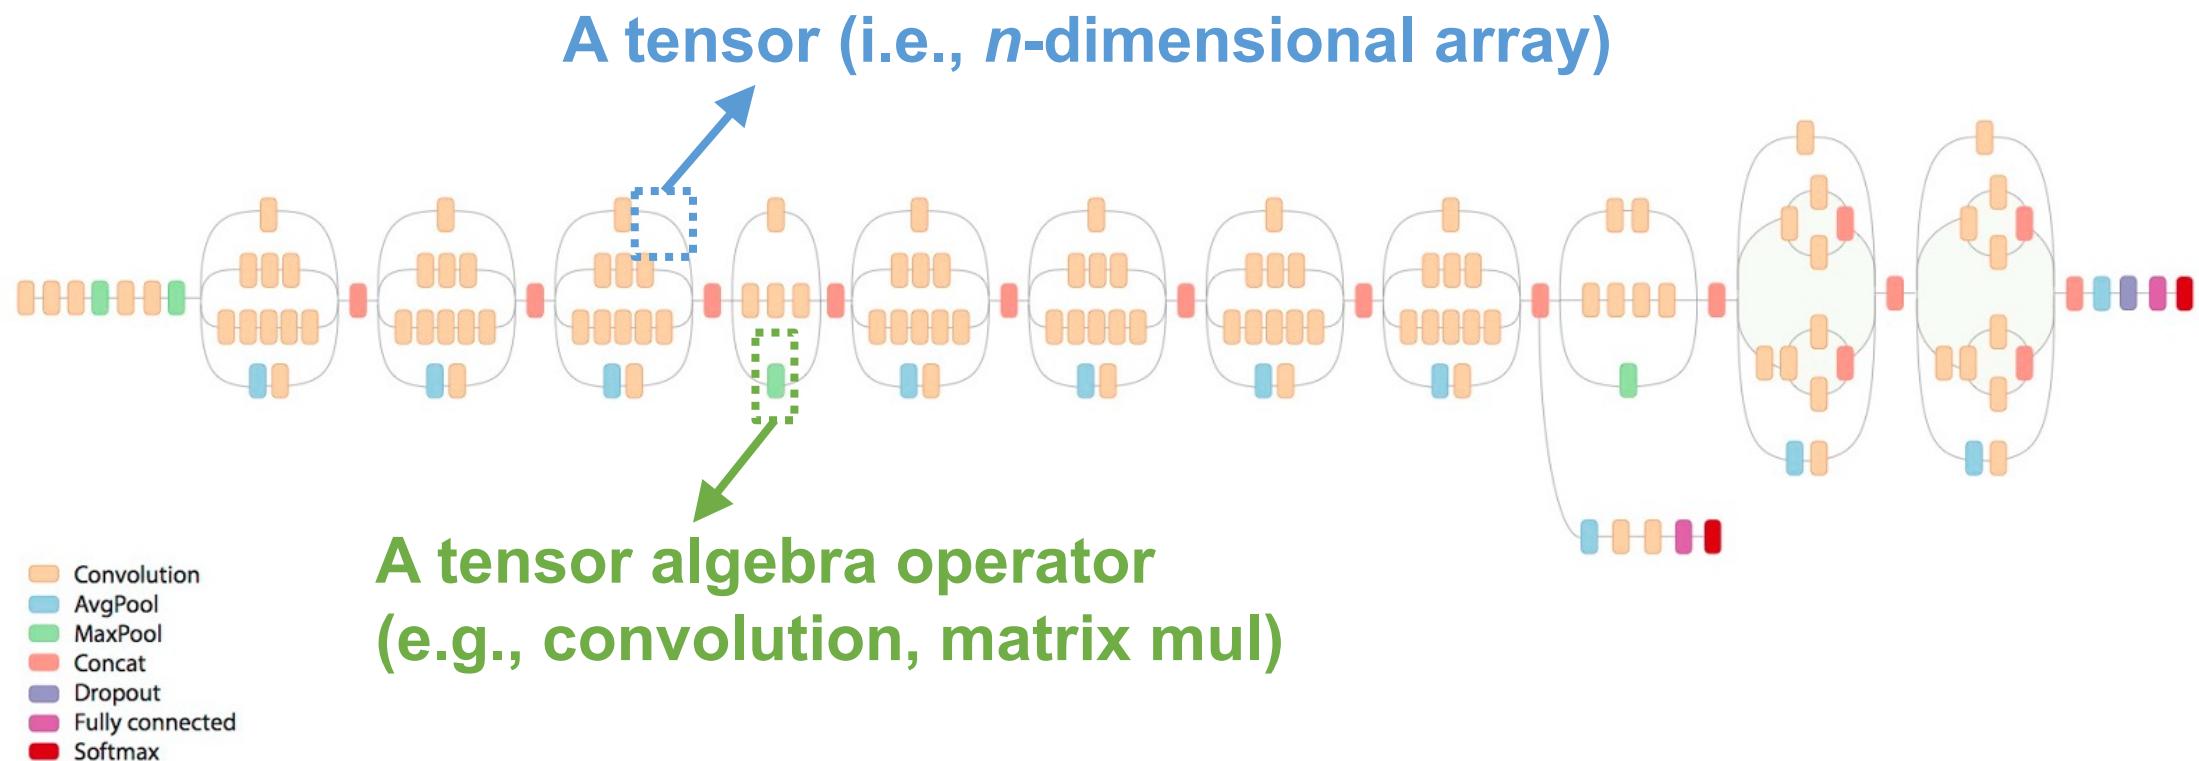
1. Halide: for image processing
2. TVM: for deep learning

The Success of Machine Learning Today



Deep Neural Network

- Collection of simple trainable mathematical units that work together to solve complicated tasks



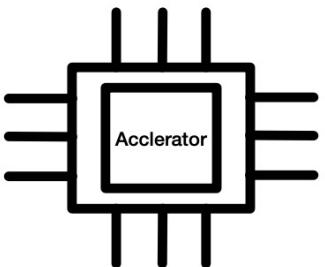
TVM: A Learning-based Compiler for Deep Learning



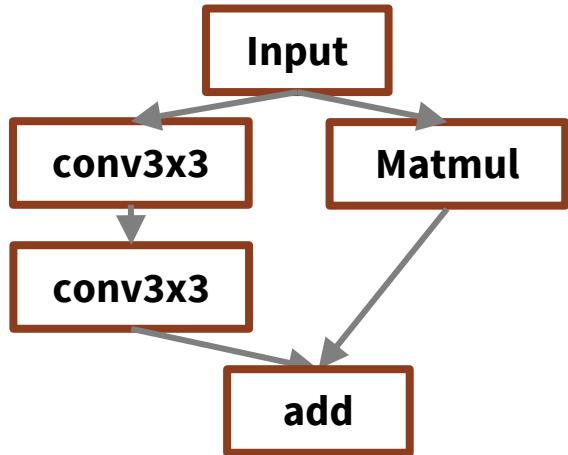
Explosion of models and frameworks

Goal: efficiently deploy deep learning on modern hardware platforms

Explosion of hardware backends



Existing Approach: Engineer Optimized Tensor Operators



Matmul: Operator Specification

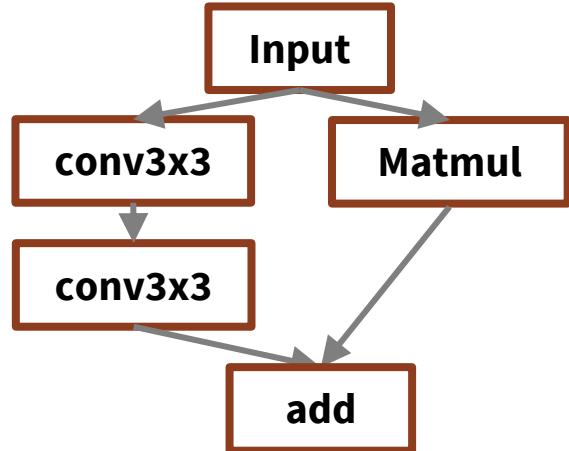
```
| C = tvm.compute((m, n),  
|   lambda y, x: tvm.sum(A[k, y] * B[k, x], axis=k))
```



Vanilla Code

```
for y in range(1024):  
    for x in range(1024):  
        C[y][x] = 0  
        for k in range(1024):  
            C[y][x] += A[k][y] * B[k][x]
```

Existing Approach: Engineer Optimized Tensor Operators



Matmul: Operator Specification

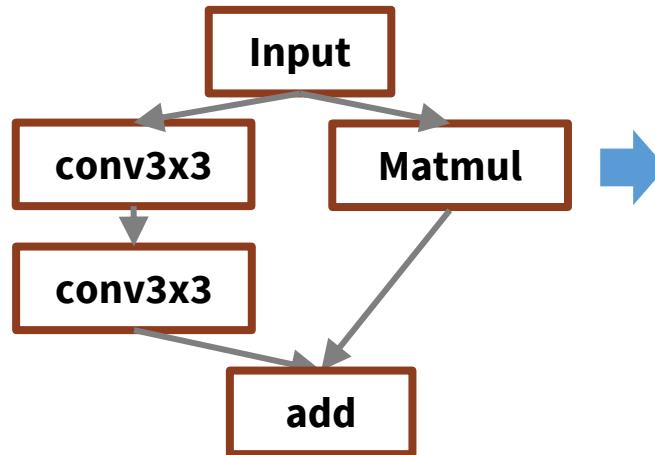
```
C = tvm.compute((m, n),  
    lambda y, x: tvm.sum(A[k, y] * B[k, x], axis=k))
```



Loop Tiling for Locality

```
for yo in range(128):  
    for xo in range(128):  
        C[yo*8:yo*8+8][xo*8:xo*8+8] = 0  
        for ko in range(128):  
            for yi in range(8):  
                for xi in range(8):  
                    for ki in range(8):  
                        C[yo*8+yi][xo*8+xi] +=  
                            A[ko*8+ki][yo*8+yi] * B[ko*8+ki][xo*8+xi]
```

Existing Approach: Engineer Optimized Tensor Operators



Matmul: Operator Specification

```
C = tvm.compute((m, n),  
    lambda y, x: tvm.sum(A[k, y] * B[k, x], axis=k))
```

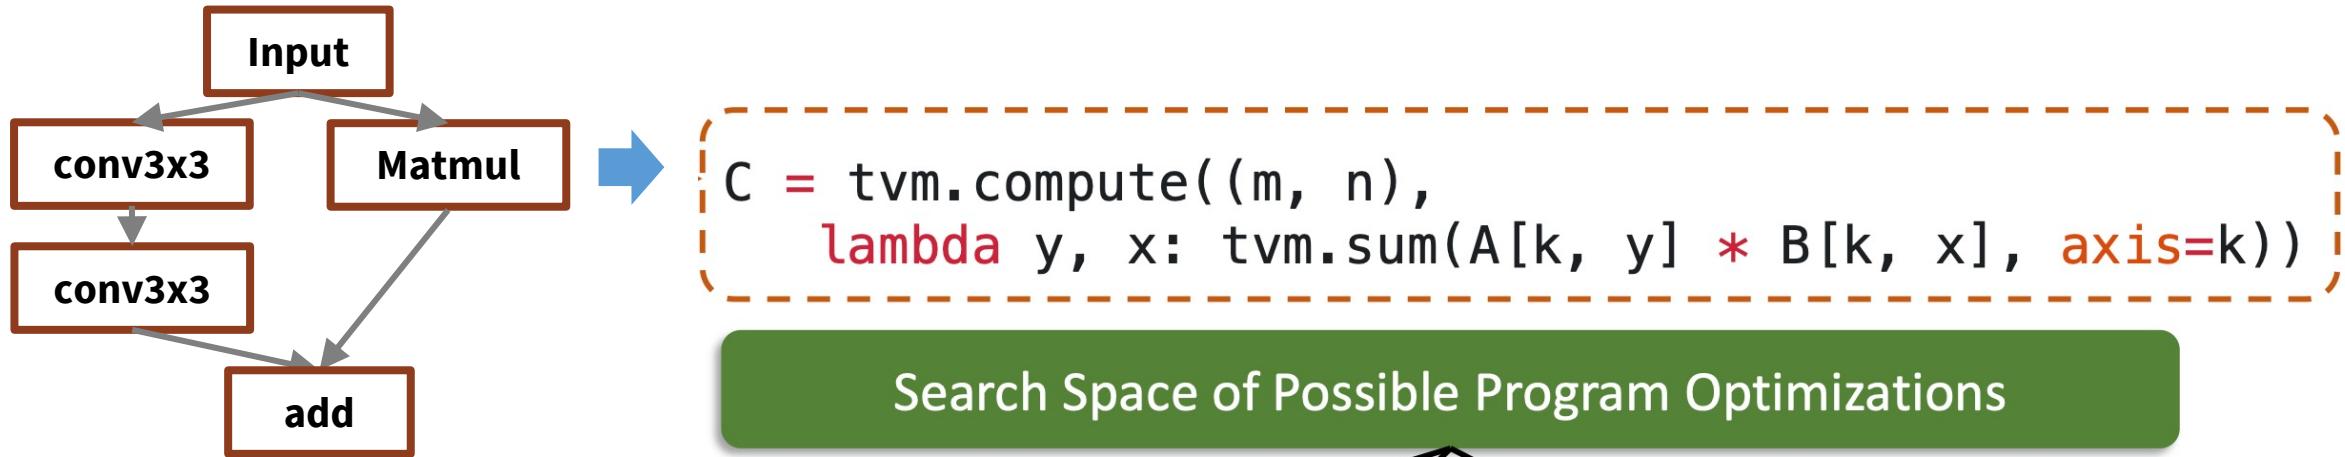


Map to Accelerators

```
inp_buffer AL[8][8], BL[8][8]  
acc_buffer CL[8][8]  
for yo in range(128):  
    for xo in range(128):  
        vdla.fill_zero(CL)  
        for ko in range(128):  
            vdla.dma_copy2d(AL, A[ko*8:ko*8+8][yo*8:yo*8+8])  
            vdla.dma_copy2d(BL, B[ko*8:ko*8+8][xo*8:xo*8+8])  
            vdla.fused_gemm8x8_add(CL, AL, BL)  
            vdla.dma_copy2d(C[yo*8:yo*8+8], xo*8:xo*8+8], CL)
```

Human exploration of optimized code

Challenge: Billions of Possible Optimization Choices in the Search Space



Low-level Program Variants

```
inp_buffer AL[8][8], BL[8][8]
acc_buffer CL[8][8]
for yo in range(128):
    for xo in range(128):
        vdla.fill_zero(CL)
        for ko in range(128):
            vdla.dma_copy2d(AL, A[ko*8:ko*8+8][yo*8:yo*8+8])
            vdla.dma_copy2d(BL, B[ko*8:ko*8+8][yo*8:yo*8+8])
            vdla.fused_gemm8x8_add(CL, AL, BL)
            vdla.dma_copy2d(C[yo*8:yo*8+8], CL, xo*8:xo*8+8)
```

```
for yo in range(128):
    for xo in range(128):
        C[yo*8:yo*8+8][xo*8:xo*8+8] = 0
        for ko in range(128):
            for yi in range(8):
                for xi in range(8):
                    for ki in range(8):
                        C[yo*8+yi][xo*8+xi] +=
                            A[ko*8+ki][yo*8+yi] * B[ko*8+ki][xo*8+xi]
```

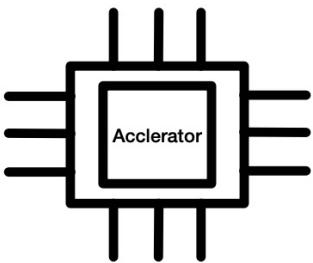
```
for y in range(1024):
    for x in range(1024):
        C[y][x] = 0
        for k in range(1024):
            C[y][x] += A[k][y] * B[k][x]
```

TVM: Learning-based Compiler for Deep Learning

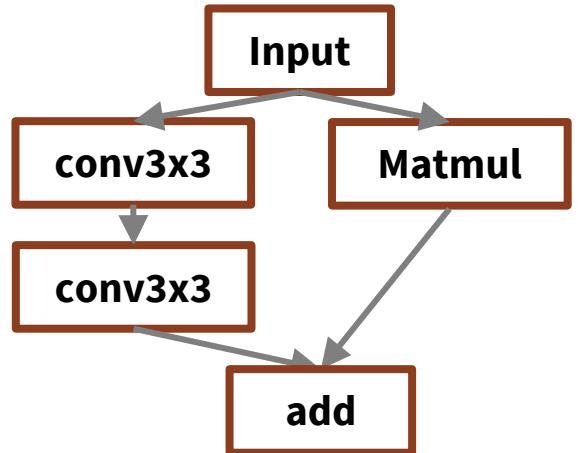


Hardware-aware Search Space of Optimized Tensor Programs

Machine Learning based Program Optimizer



Hardware-aware Search Space



≈ Halide's algorithm

Tensor Expression Language (Specification)

```
C = tvm.compute((m, n),  
    lambda y, x: tvm.sum(A[k, y] * B[k, x], axis=k))
```

Define search space of hardware aware mappings from expression to hardware program

Based on Halide's compute/schedule separation

Hardware



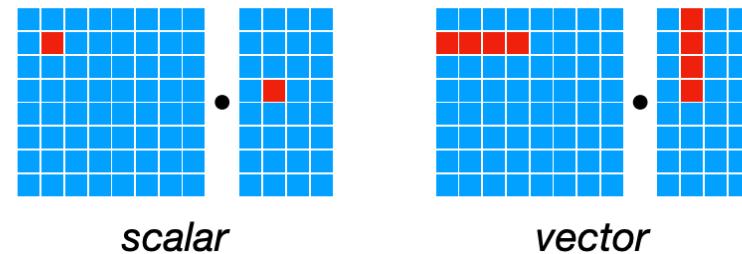
Hardware-aware Search Space

Reuse primitives from Halide

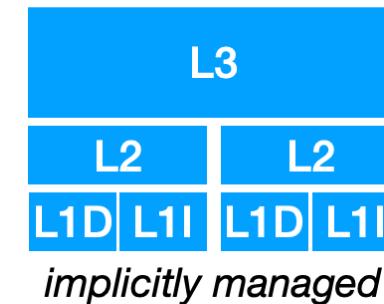
CPUs



Compute Primitives



Memory Subsystem



Loop
Transformations

Cache
Locality

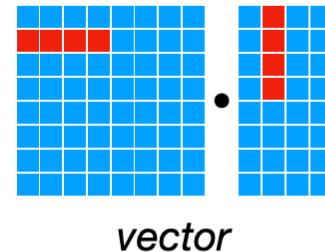
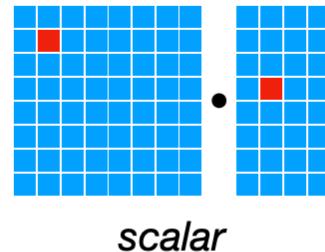
Vectorization

Hardware-aware Search Space

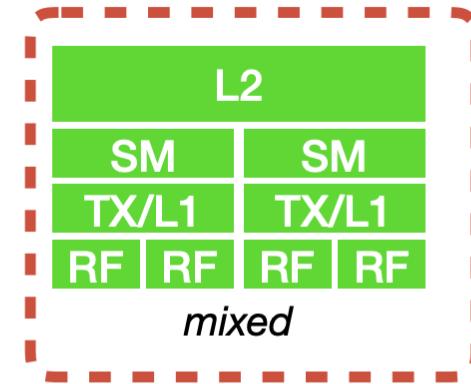
GPUs



Compute Primitives



Memory Subsystem



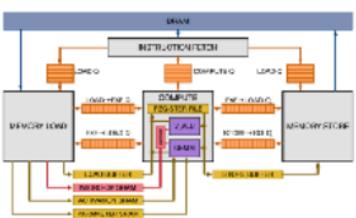
Shared memory among
compute cores

Use of Shared
Memory

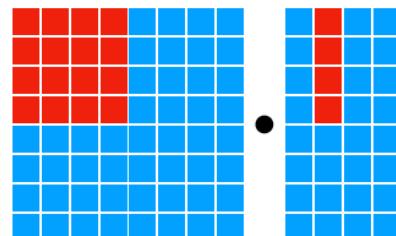
Thread
Cooperation

Hardware-aware Search Space

TPU-like Specialized Accelerators



Compute Primitives



tensor

Tensorization

Memory Subsystem

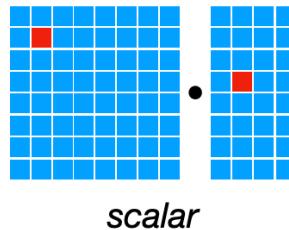
Unified Buffer

FIFO
Acc

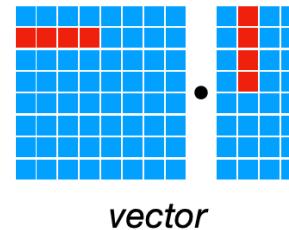
explicitly managed

Tensorization Challenge

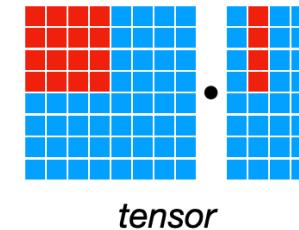
Compute primitives



scalar



vector



tensor

Hardware designer:
declare tensor instruction interface
with Tensor Expression

```
w, x = t.placeholder((8, 8)), t.placeholder((8, 8))
k = t.reduce_axis(0, 8)
y = t.compute(8, 8, lambda i, j:
    t.sum(w[i, k] * x[j, k], axis=k))

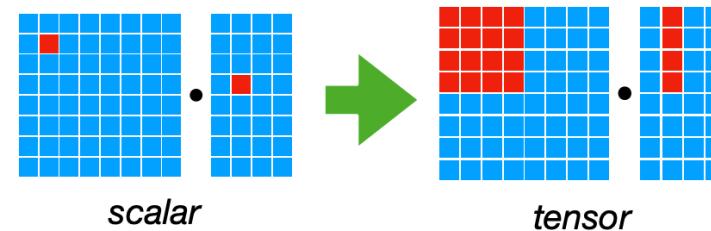
def gemm_intrinsic_lower(inputs, outputs):
    ww_ptr = inputs[0].access_ptr("r")
    xx_ptr = inputs[1].access_ptr("r")
    zz_ptr = outputs[0].access_ptr("w")
    compute = t.hardware_intrinsic("gemm8x8", ww_ptr, xx_ptr, zz_ptr)
    reset = t.hardware_intrinsic("fill_zero", zz_ptr)
    update = t.hardware_intrinsic("fuse_gemm8x8_add", ww_ptr, xx_ptr, zz_ptr)
    return compute, reset, update

gemm8x8 = t.decl_tensor_intrinsic(y.op, gemm_intrinsic_lower)
```

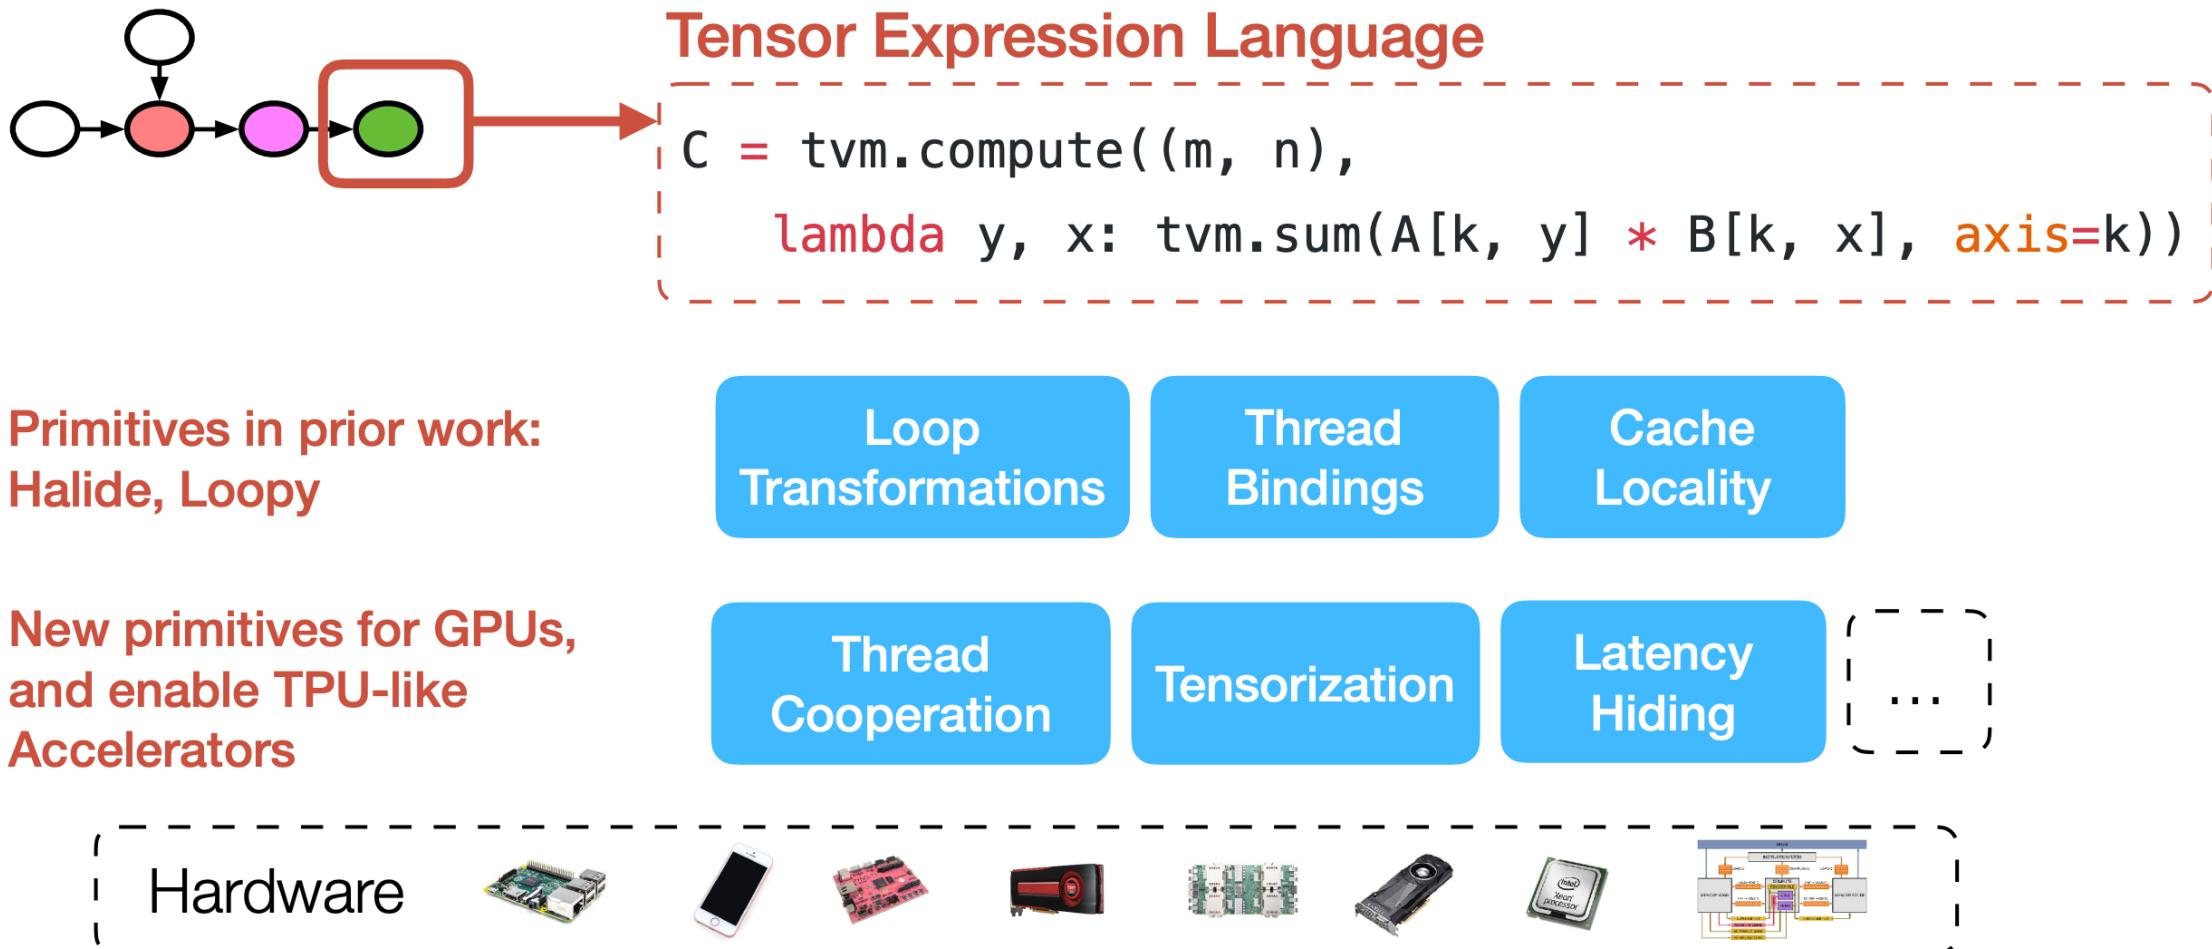
declare behavior

lowering rule to generate hardware intrinsics to carry out the computation

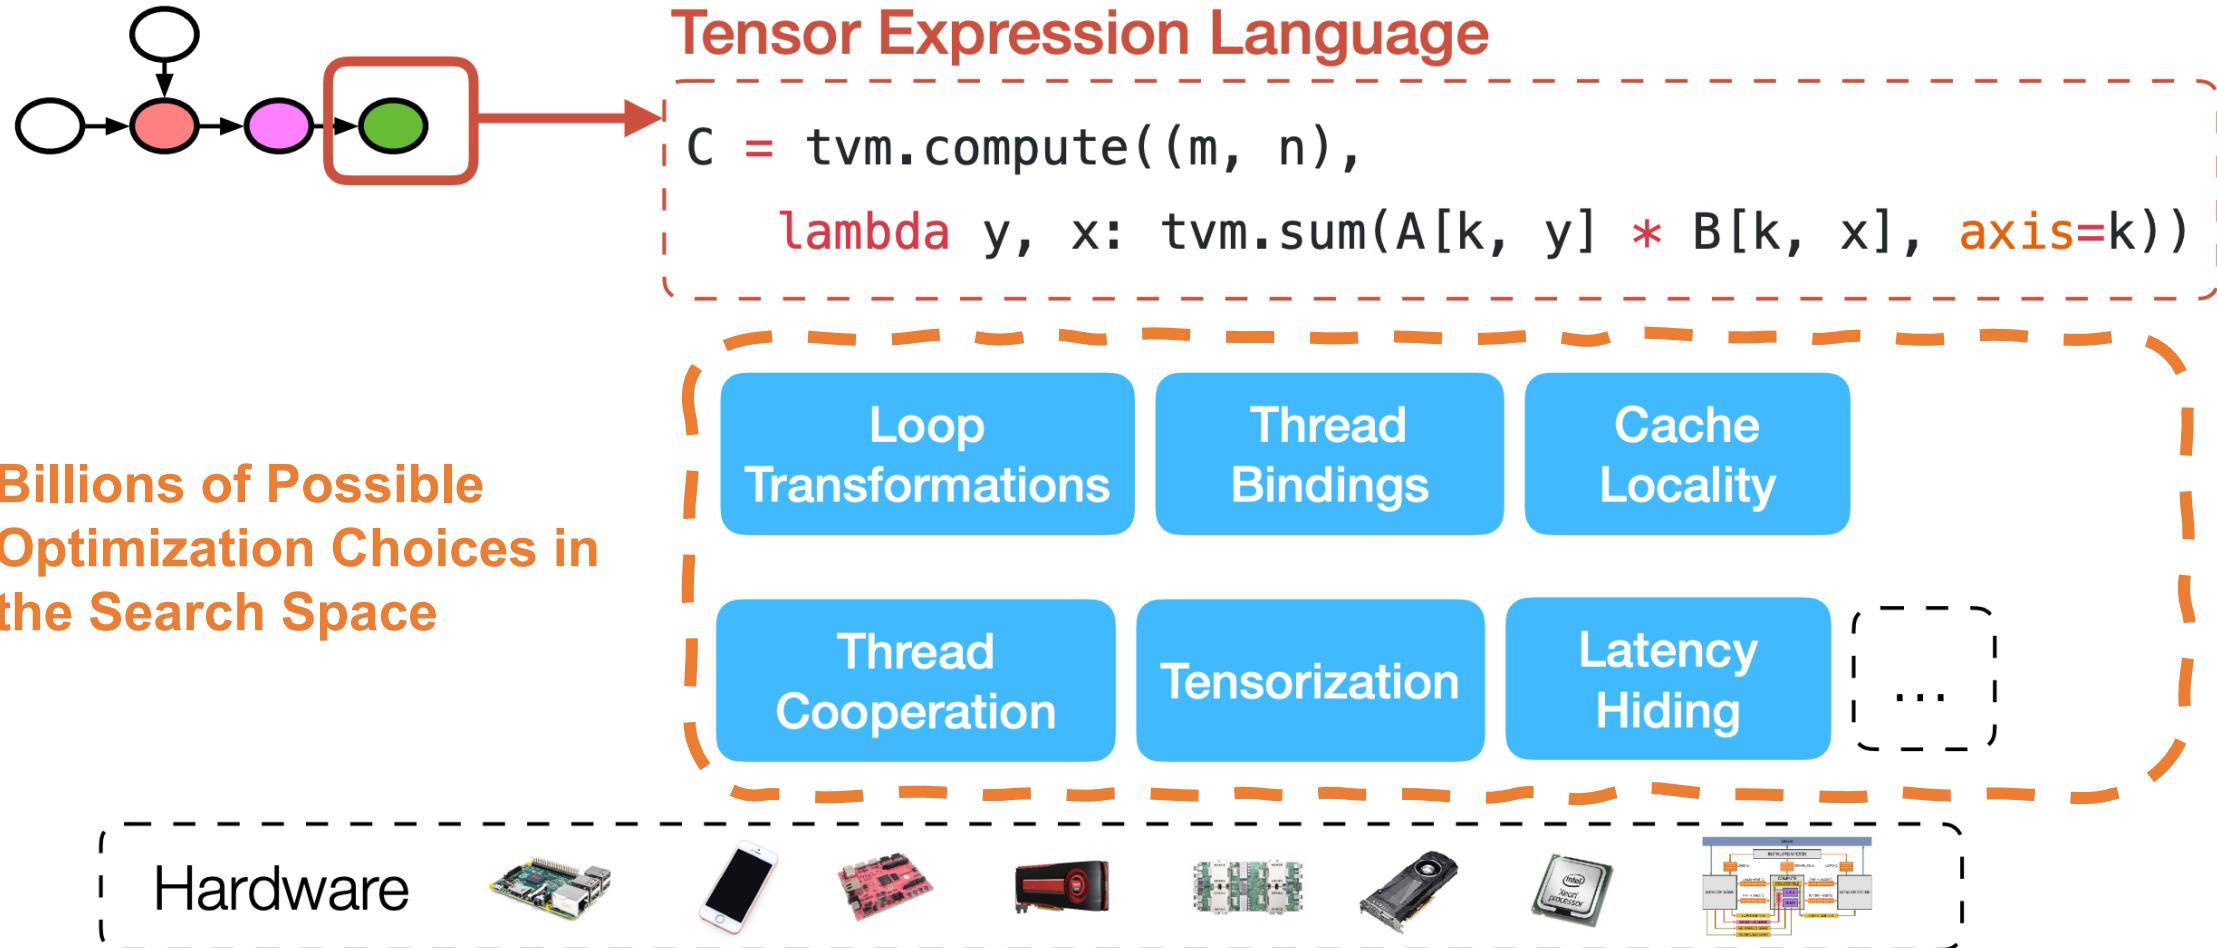
Tensorize:
transform program to use tensor instructions



Hardware-aware Search Space



Hardware-aware Search Space

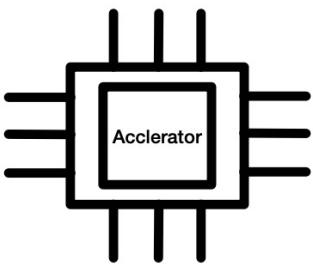


TVM: Learning-based Compiler for Deep Learning

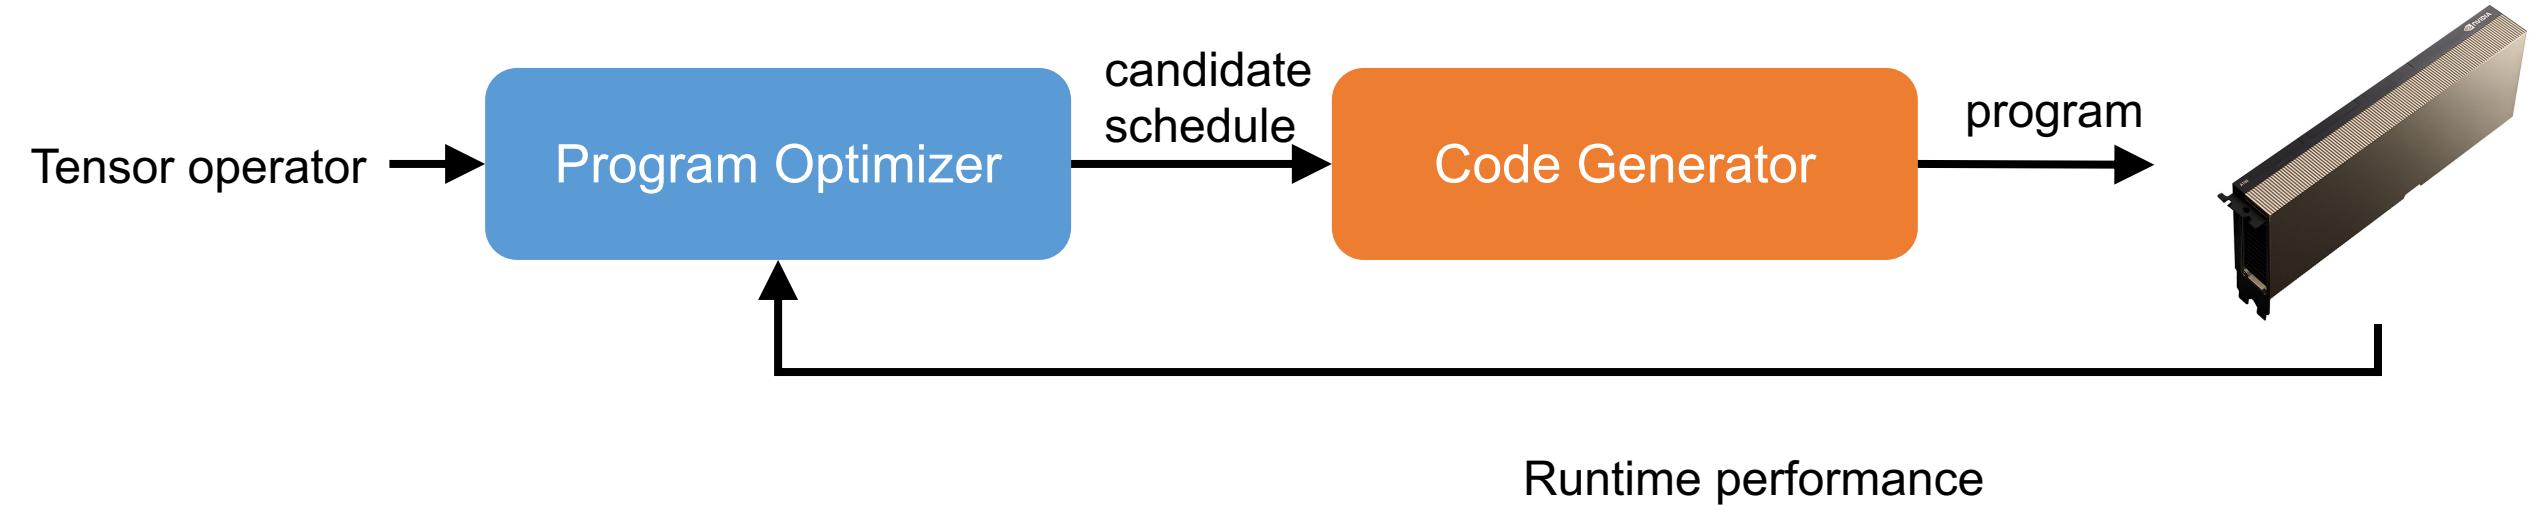


Hardware-aware Search Space of Optimized Tensor Programs

Learning based Program Optimizer

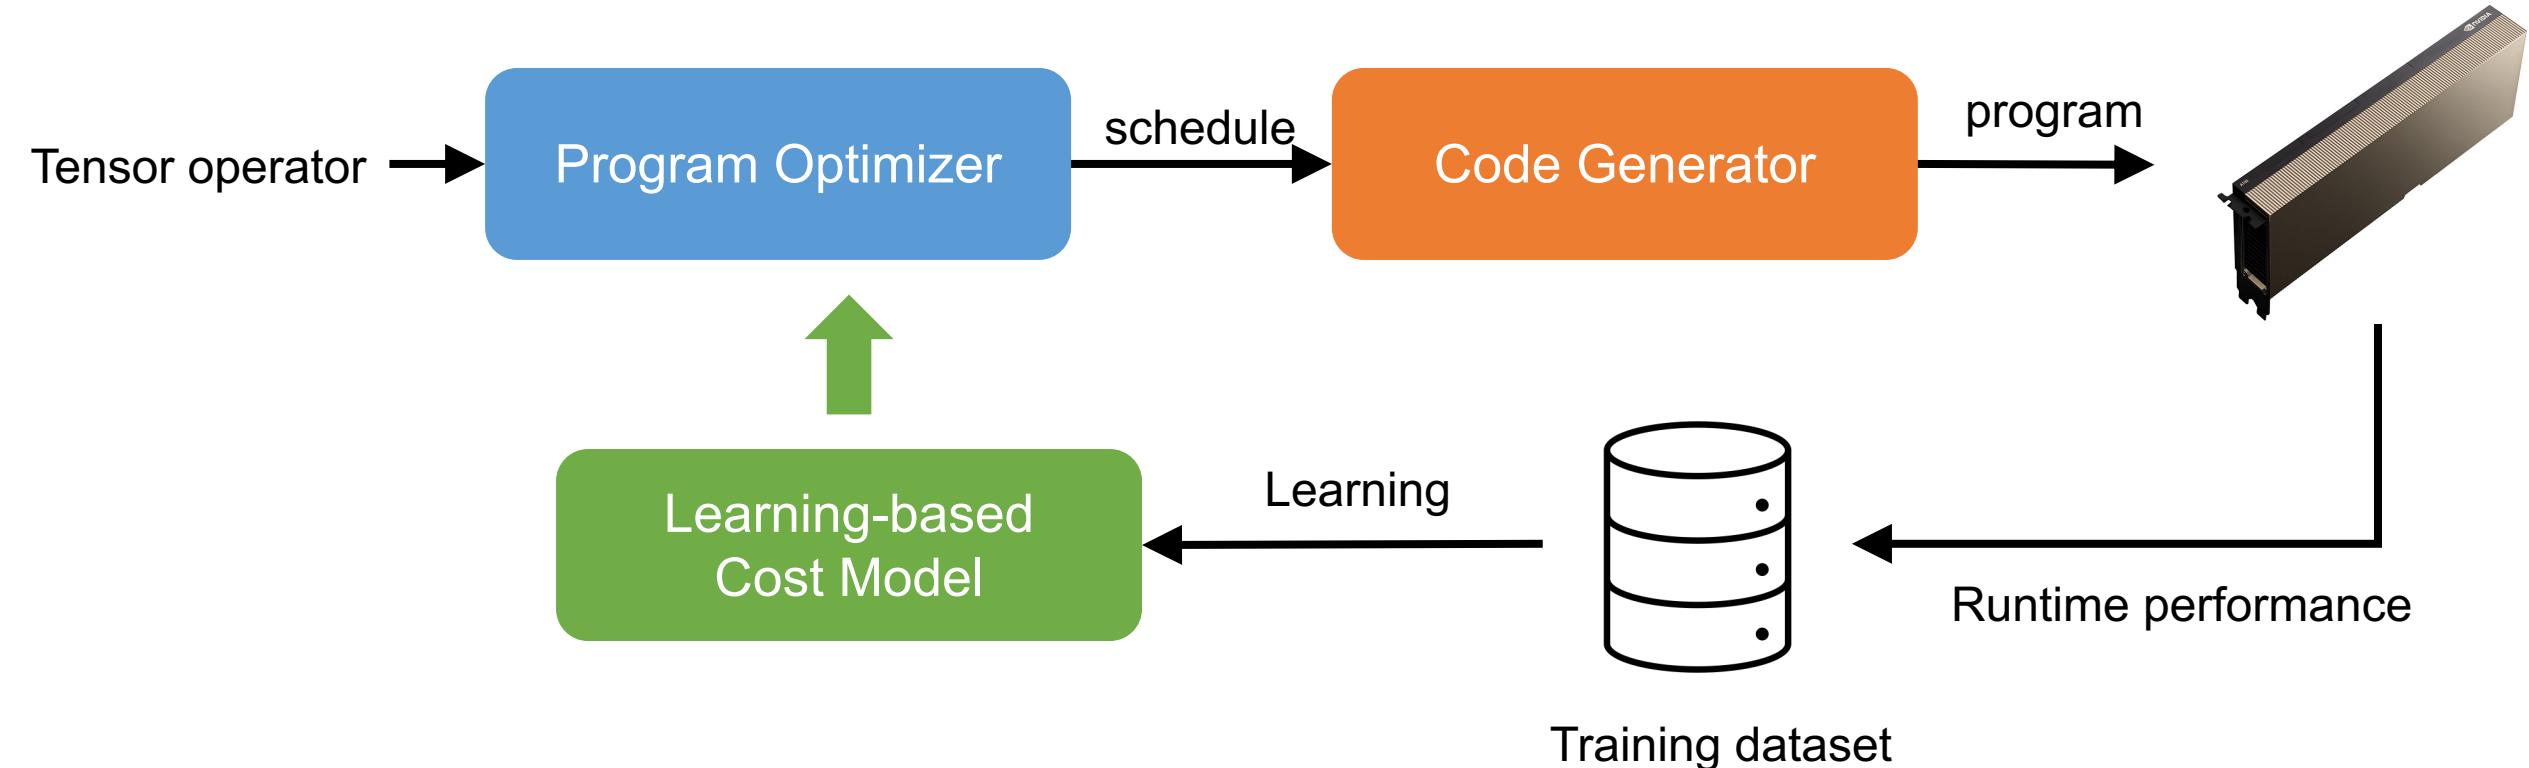


Learning-based Program Optimizer



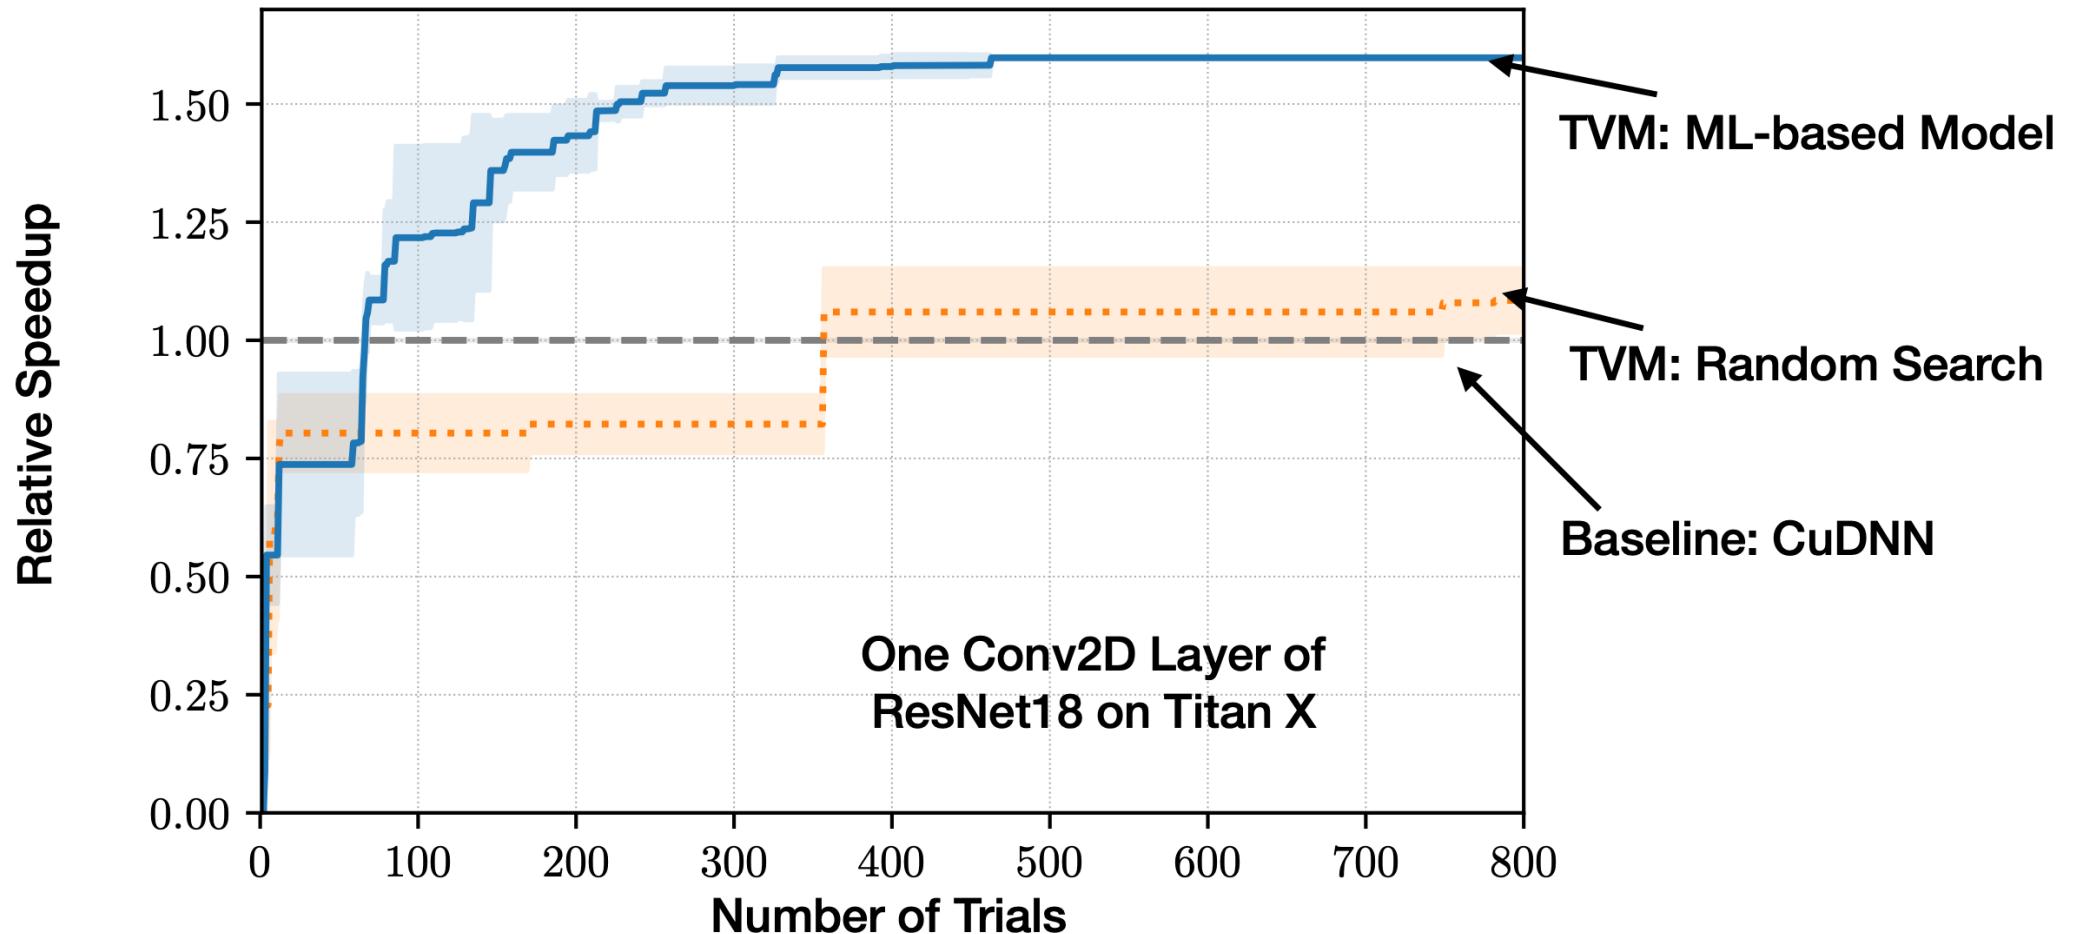
Issue: high experiment cost, each trial takes seconds

Learning-based Program Optimizer

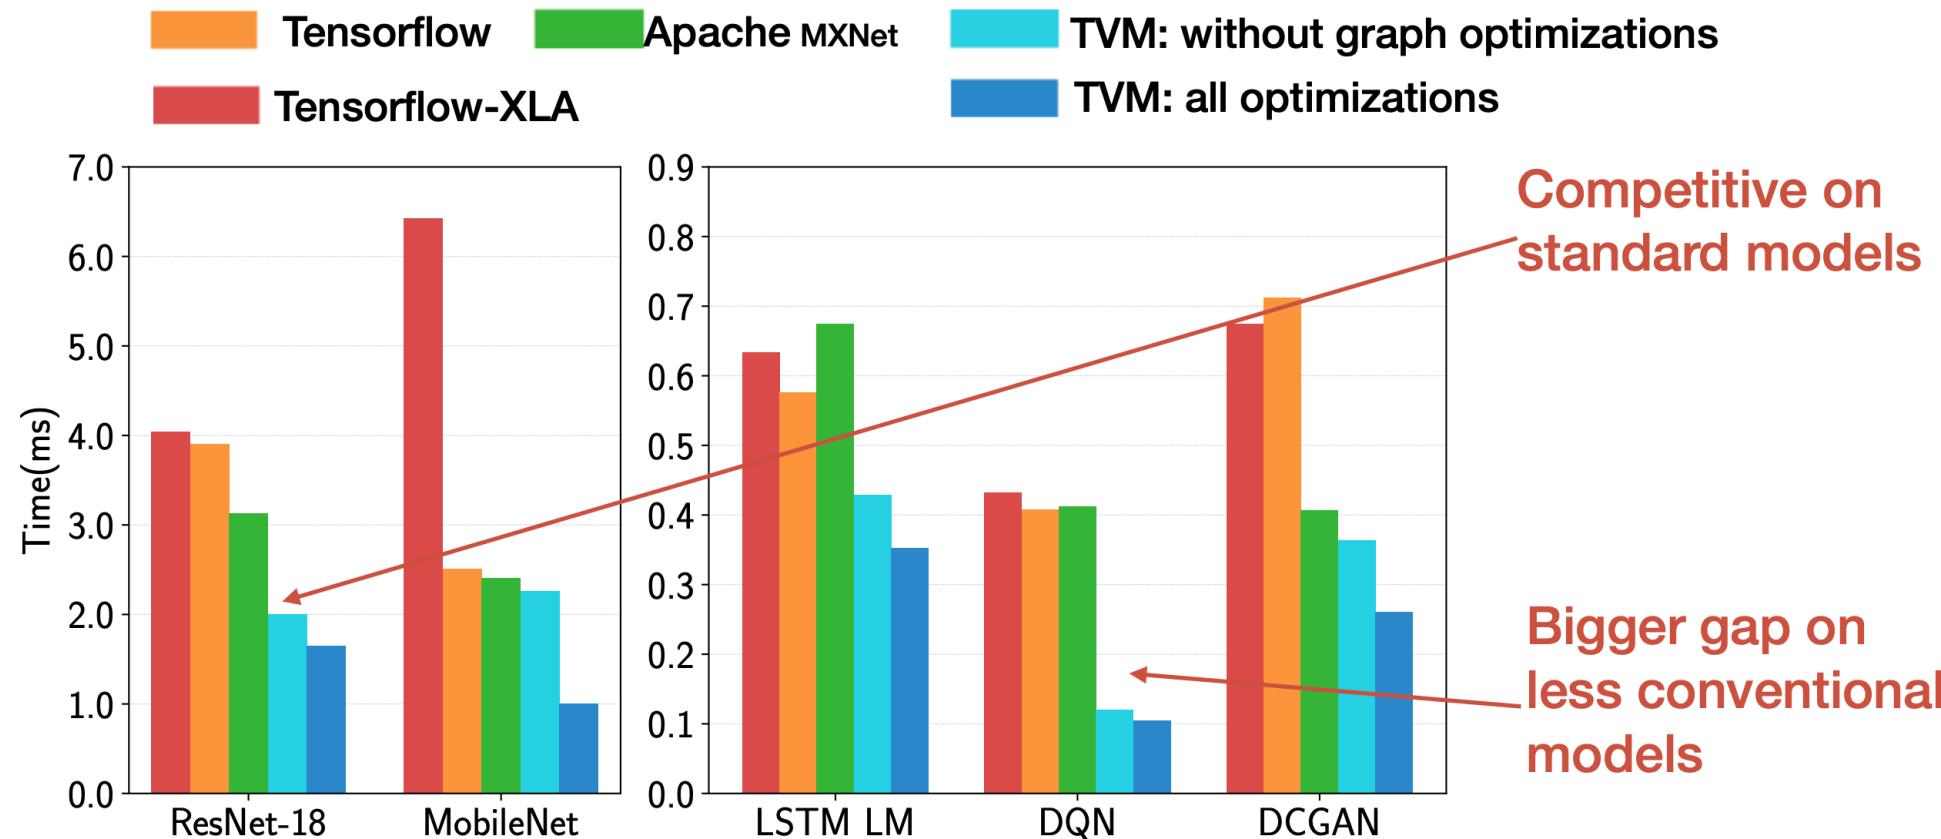


Adapt to hardware by learning, make prediction in milliseconds

Efficient ML-based Cost Model



End-to-end Inference Performance



Discussion: Halide and TVM

- What are the similarities?
- What are the key differences?

Summary

- Modern machines are parallel and heterogeneous
 - Only way to increase compute capability in energy-constrained world
- Most software uses small fraction of peak capability of machine
 - Challenging to tune programs to these machines
 - Tuning efforts not portable across machines
- DSLs trade-off **generality** to achieve **productivity, performance, portability**
 - Case studies: Halide, TVM